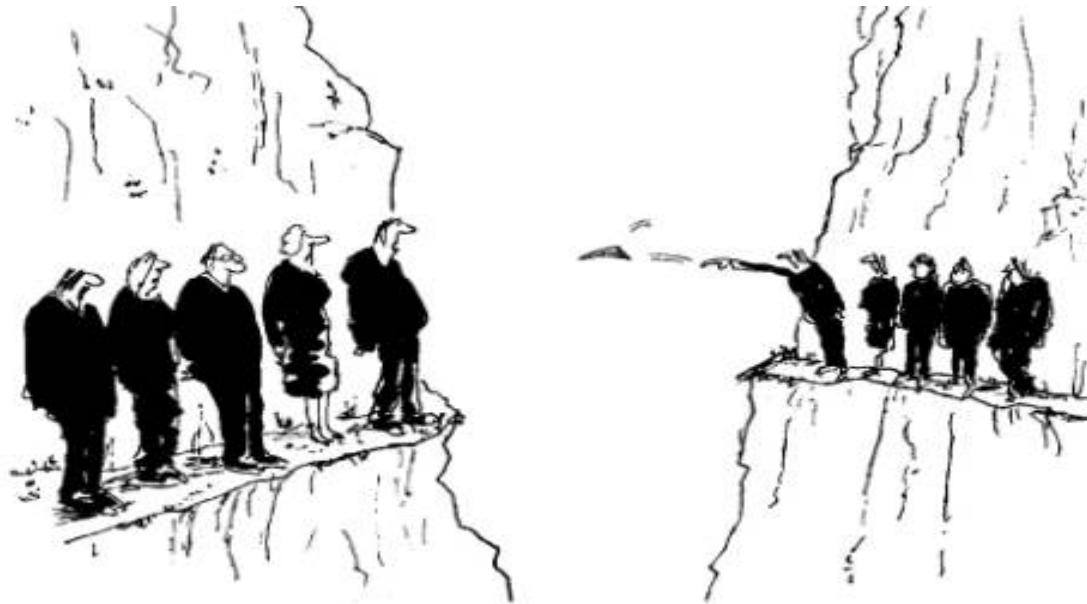


Complexity, Tradeoffs, Architecture, and the Evolution of DevOps



Couldn't we communicate better if we built a bridge?

David Meyer

CTO and Chief Scientist, Brocade

Director, Advanced Technology Center, University of Oregon

dmm@{brocade.com,uoregon.edu,1-4-5.net,...}

<http://www.1-4-5.net/~dmm/talks/bridges.pdf>

Agenda

- Goals for this Talk
- What is Complexity and Why is it Hidden
 - Robustness, Fragility, and Complexity
- Tradeoffs?
- The Architecture of Complex Systems
 - Universal Architectural Principles
- The Evolution DevOps
- A Few Conclusions and Q&A if we have time

Danger Will Robinson!!!



*This talk might be controversial/provocative
(and perhaps a bit “sciencey”)*

Standard Disclaimer

Goals for this Talk

- Characterize the essential features of *complexity*
 - in both technological and biological systems
- Understand the fundamental *tradeoffs*
 - that are made in these systems
- Understand the *universal architectural features*
 - found in both technological and biological networks
- *Evolving DevOps*
 - means bridging the DevOps and theory communities

Goals for this Talk

Said another way:

To open up our thinking about what the essential architectural features of our network are, how these features combine to provide robustness (and its dual, fragility), and how the universal architectural features that we find in both technological and biological networks effect Internet robustness, scalability and evolvability.

Agenda

- Goals for this Talk
- What is Complexity and Why is it Hidden
 - Robustness, Fragility, and Complexity
- Tradeoffs?
- The Architecture of Complex Systems
 - Universal Architectural Principles
- The Evolution of DevOps
- A Few Conclusions and Q&A if we have time

What is Complexity (and why is it Hidden)?

- So what is Complexity?
 - Complexity is hidden *structure* that arises in systems
 - Its purpose is to create *robustness* to environmental and component uncertainty
- Why hidden?
 - Anti-lock/anti-skid brakes, packet loss, OS kernels, power grids, SDN controllers, lunar landing systems, ...
 - You don't notice they are there...until they fail
 - often catastrophically
 - The hidden nature of complexity is a fundamental property of systems
 - and derives from universal architectural principles of complex systems
 - we will see examples of this in just a minute
- Ok, then what is Robustness?

Robustness is a Generalized System Feature

- **Scalability** is *robustness* to changes to the size and complexity of a system as a whole
- **Evolvability** is *robustness* of lineages to changes on various (usually long) time scales
- Other system features cast as robustness
 - **Reliability** is *robustness* to component failures
 - **Efficiency** is *robustness* to resource scarcity
 - **Modularity** is *robustness* to component rearrangements
- Of course, these are the features we're seeking from the network

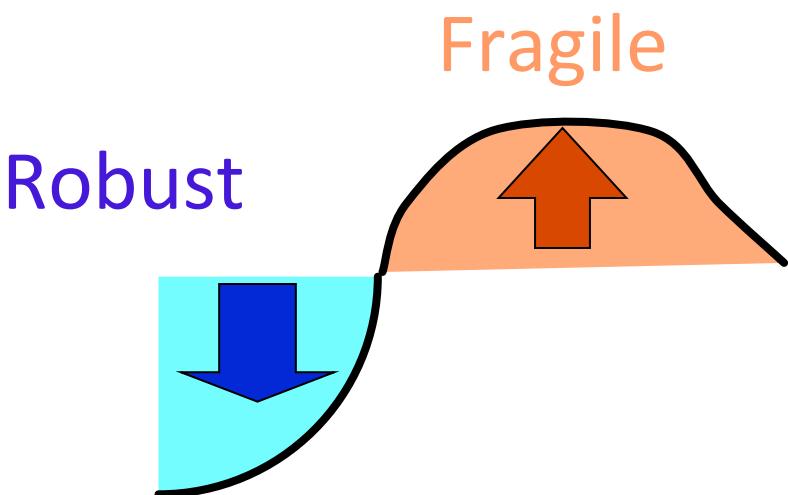
A Bit More Formally

- **Robustness** is the preservation of a certain property in the presence of uncertainty in components or the environment
 - Obviously a core Internet design principle
 - Systems Biology: Biological systems are designed such that their important functions are insensitive to the naturally occurring variations in their parameters.
 - Limits the number of designs that can actually work in the real environment
 - Exact adaptation in bacteria chemotaxis
- **Fragility** is the opposite of robustness
 - Another way to think about fragility
 - Technical: You are fragile if you depend on 2nd order effects (acceleration) and the “harm” curve is concave
 - A little more on this in the next few slides...
- A system can have a *property* that is *robust* to one set of perturbations and yet *fragile* for a *different property* and/or perturbation → the system is **Robust Yet Fragile**
 - Or the system may collapse if it experiences perturbations above a certain threshold (K-fragile)
- For example, a possible **RYF tradeoff** is that a system with high efficiency (i.e., using minimal system resources) might be unreliable (i.e., fragile to component failure) or hard to evolve
 - VRRP, ISSU, HA, TE, ...
 - Complexity/Robustness Spirals
 - Implications for Carrier Grade components?

Robust Yet Fragile?

[a system] can have
[a property] that is **robust** to
[a set of perturbations]

Yet be **fragile** for
[a different property]
Or [a different perturbation]



Recent results suggest that the RYF tradeoff is a hard tradeoff that cannot be overcome¹

This is profound: If you create robustness somewhere you will create fragility somewhere else...

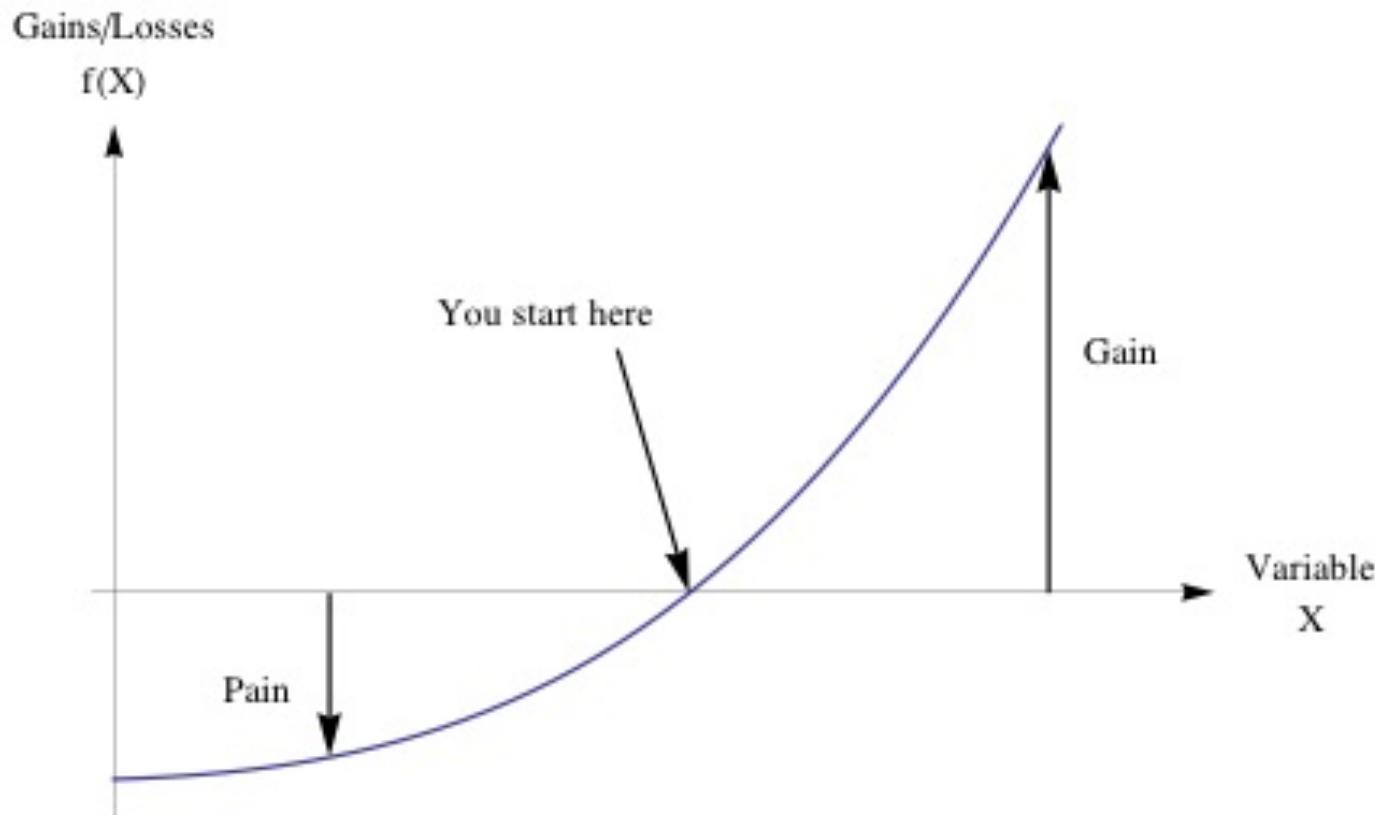
DevOps (along with most engineering disciplines) does not currently account for this effect

Harm Function: Concave → Fragile, Convex → Robust

¹ See Marie E. Csete and John C. Doyle, "Reverse Engineering of Biological Complexity", <http://www.cds.caltech.edu/~doyle/wiki/images/0/05/ScienceOnlinePDF.pdf>

Another way to think about Robustness

(Convex Optionality aka No Pain, No Gain)



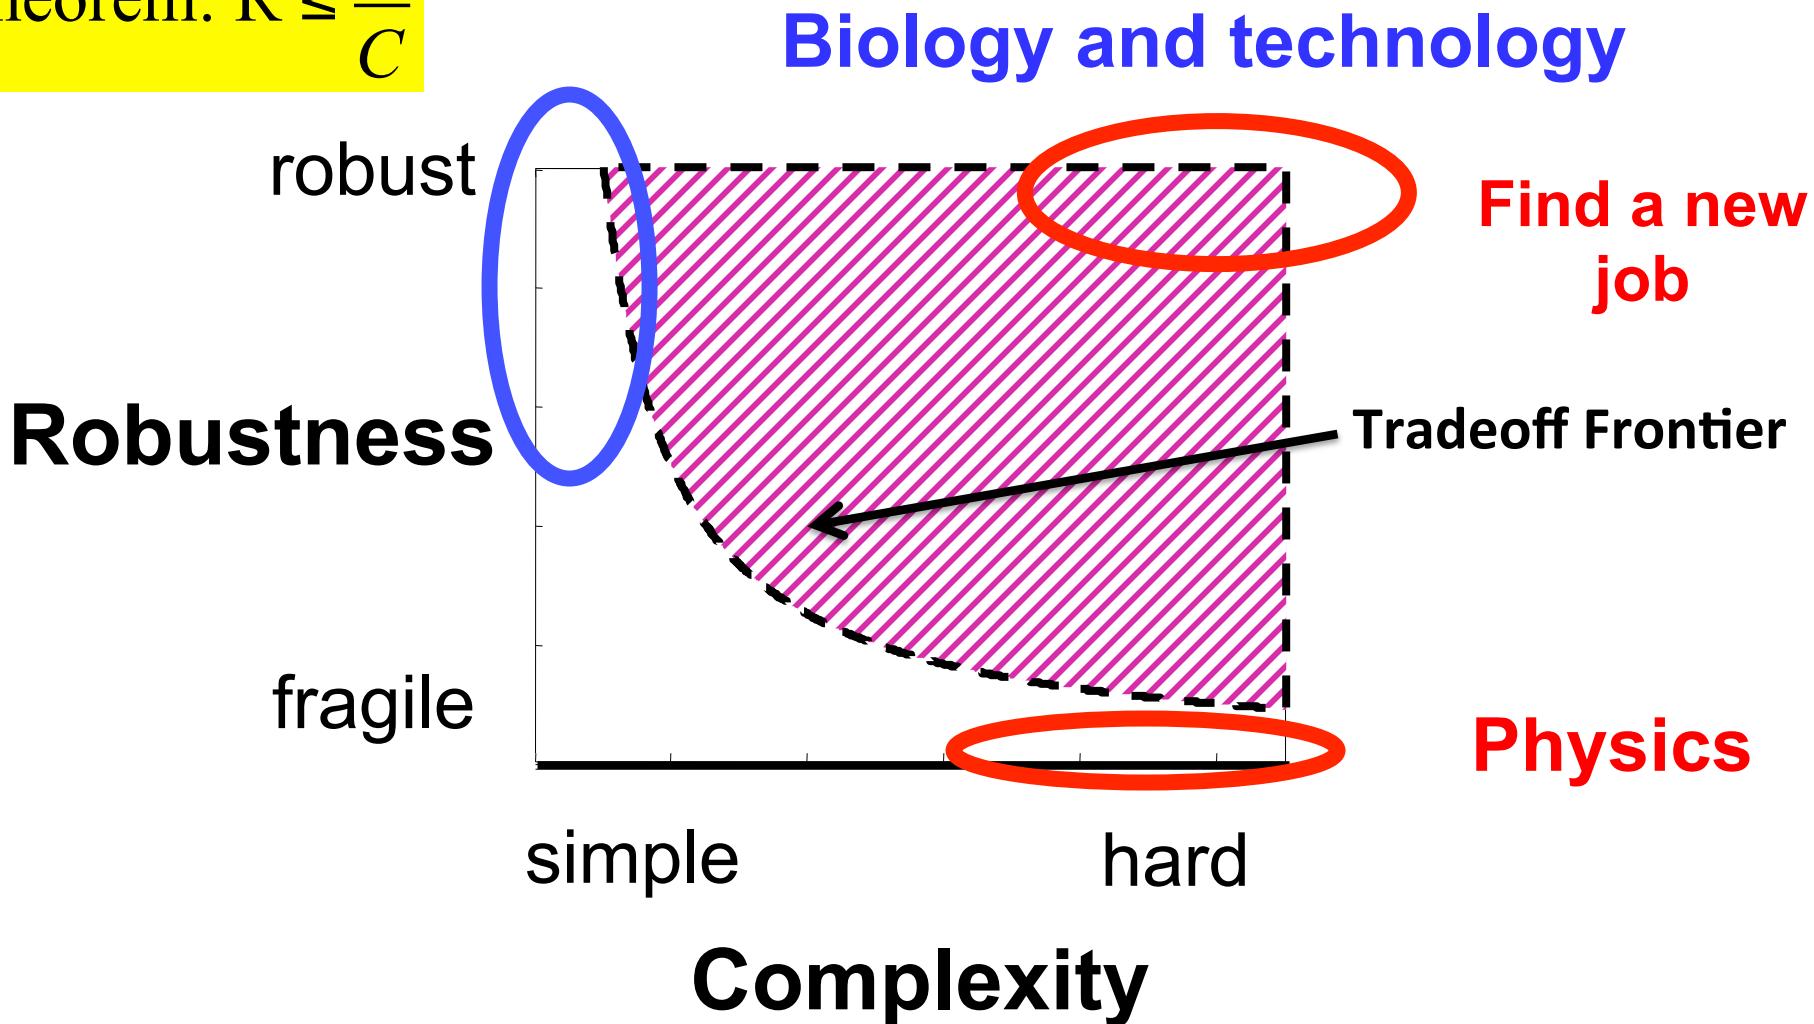
If you are hurt much less by error/variation than you stand to gain...

Interestingly, Fragility and Scaling are Related

- A bit of a formal description of fragility
 - Let z be some stress level, p some property, and
 - Let $H(p,z)$ be the (negative valued) harm function
 - Then for the fragile the following must hold
 - $H(p,nz) < nH(p,z)$ for $0 < nz < K$
- → *A big event hurts non-linearly more than the sum of small events*
- For example, a coffee cup on a table suffers non-linearly more from large deviations ($H(p, nz)$) than from the cumulative effect of smaller events ($nH(p,z)$)
 - So the cup is damaged far more by *tail events* than those within a few σ 's of the mean
 - Sensitivity to tail events → RYF
 - Too theoretical? Perhaps, but consider: ARP storms, micro-loops, congestion collapse, AS 7007, ...
 - BTW, nature requires this property
 - Consider: jump off something 1 foot high 30 times v/s jumping off something 30 feet high once
- So when we say something scales like $O(n^2)$, what we mean is the damage to the network has constant acceleration (2) for *weird* enough n (e.g., outside say, 10 σ)
 - Again, ARP storms, congestion collapse, AS 7007, DDOS, ... → non-linear damage

Its All About (RYF) Tradeoffs

Theorem: $R \leq \frac{1}{C}$



BTW, RYF Behavior is Everywhere

Robust

- ☺ Efficient, flexible metabolism
- ☺ Complex development
- ☺ Immune systems
- ☺ Regeneration & renewal
- 📄 Complex societies
- 📄 Advanced Technologies

Yet Fragile

- ☹ Obesity and diabetes
- ☹ Rich microbe ecosystem
- ☹ Inflammation, Auto-Im.
- ☹ Cancer
- ☠ Epidemics, war, ...
- 💣 Catastrophic failures

- “Evolved” mechanisms for robustness *allow for*, even *facilitate*, novel, severe fragilities elsewhere. That is, they are RYF-Complex
- Often involving hijacking/exploiting the same mechanism
 - We’ve certainly seen this in the Internet space (consider DDOS of various varieties)
- These are hard constraints (i.e., RYF behavior is conserved)

Robust

- 😊 Metabolism
- 😊 Regeneration
- 😊 Healing w/o

Fragile

- 😢 Obesity, diabetes

- 😢 Fat accumulation
- 😢 Insulin resistance
- 😢 Proliferation
- 😢 Inflammation

Same mechanisms

- Fragility ← Hijacking, side-channel attacks
- DDoS, reflection attacks
- Of mechanisms → robustness
- Complexities → robust/fragile tradeoffs
- Mechanisms → constraints (“conservation laws”)

Both

Accident or necessity?

Summary: Understanding RYF is *The Challenge*

- It turns out that managing/understanding RYF behavior is ***the most essential challenge*** in technology, society, politics, ecosystems, medicine, etc. This means...
 - Understanding *Universal Architectural Principles*
 - Look ahead: Layering, Bowties/Hourglasses, Constraints that Deconstrain
 - Managing spiraling complexity/fragility
 - Not predicting what is likely or typical
 - But rather understanding what is catastrophic
 - or in Taleb's terminology, that which is fat tailed
- BTW, it is much easier to create the robust features than it is to prevent the fragilities
 - And as I mentioned, there are poorly understood “conservation laws” at work¹
- Bottom Line
 - ***Understanding RYF behavior means understanding architecture and the hidden nature of complexity***

¹ See Marie E. Csete and John C. Doyle, “Reverse Engineering of Biological Complexity”,
<http://www.cds.caltech.edu/~doyle/wiki/images/0/05/ScienceOnlinePDF.pdf>

Agenda

- ~~Goals for this Talk~~
- ~~What is Complexity and Why is it Hidden~~
 - ~~Robustness, Fragility, and Complexity~~
- Tradeoffs?
- The Architecture of Complex Systems
 - Universal Architectural Principles
- The Evolution of DevOps
- A Few Conclusions and Q&A if we have time

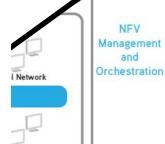
What is our Architecture (and what tradeoffs are we making)

Higher Grade,

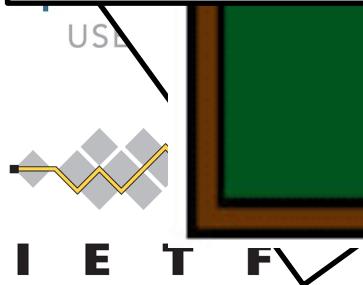


openstack

Open source is good for me. I will fully embrace it.
Open source is good for me. I will fully embrace it.



So what are the fundamental tradeoffs that we are making, and is there a more general way to think about them? But first...



First, what tradeoffs do we see everyday?

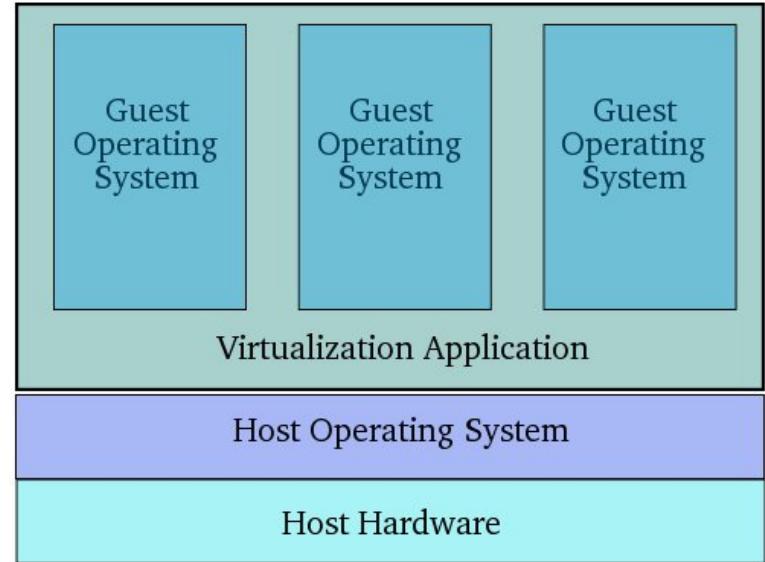
```
00 00 00 00 00 00 00 03 00 00 00 00 00 00 00 00 .....  
00 00 00 14 00 00 00 00 00 00 00 07 00 00 00 ..`.....H.....  
00 00 00 00 00 00 A0 03 40 00 00 00 00 00 .....@.....@.....  
00 00 00 18 00 00 00 00 00 00 00 FE FF FF 6F .....0.....0.....  
00 00 00 00 00 00 F0 FF FF 6F 00 00 00 00 .....@.....0.....0....  
00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 t.@.....  
00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....  
00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....  
00 00 00 00 00 00 36 04 40 00 00 00 00 00 ....P.`.....6.@.....  
00 00 00 00 00 00 00 47 43 43 3A F.@.....V.@.....GCC:  
2D 31 75 62 75 6E 74 75 35 29 20 34 2E 36 2E (Ubuntu/Linaro 4.6.3-1ubuntu5) 4.6.  
68 73 74 72 74 61 62 00 2E 69 6E 74 65 72 70 3...syntab..strtab..shstrtab..interp  
67 6E 75 2E 62 75 69 6C 64 2D 69 64 00 2E 67 ..note.ABI-tag..note.gnu.build-id..g  
74 72 00 2E 67 6E 75 2E 76 65 72 73 69 6F 6E nu.hash..dynsym..dynstr..gnu.version  
2E 64 79 6E 00 2E 72 65 6C 61 2E 70 6C 74 00 ..gnu.version_r..rela.dyn..rela.plt.  
64 61 74 61 00 2E 65 68 5F 66 72 61 6D 65 5F .init..text..fini..rodata..eh_frame_  
2E 64 74 6F 72 73 00 2E 6A 63 72 00 2E 64 79 hdr..eh_frame..ctors..dtors..jcr..dy  
64 61 74 61 00 2E 62 73 73 00 2E 63 6F 6D 6D namic..got..got.plt..data..bss..comm  
00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 ent.....  
00 00 00 00 00 00 00 00 00 00 00 1B 00 00 00 .....  
00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
```

```
def add5(x):  
    return x+5  
  
def dotwrite(ast):  
    nodename = getNodeName()  
    label=symbol.sym_name.get(int(ast[0]),ast[0])  
    print '%s [%label=%s' % (nodename, label),  
    if isinstance(ast[1], str):  
        if ast[1].strip():  
            print '=' % ast[1]  
        else:  
            print '';  
    else:  
        print ''];'  
        children = []  
        for n, child in enumerate(ast[1]):  
            children.append(dotwrite(child))  
        print '%s -> {' % nodename,  
        for name in children:  
            print '%s' % name,
```

What tradeoffs are being made here?

Speed vs. flexibility?

Everyday Tradeoffs



How about here?

Speed vs. flexibility (bare metal vs. VM vs. container)?

Everyday Tradeoffs



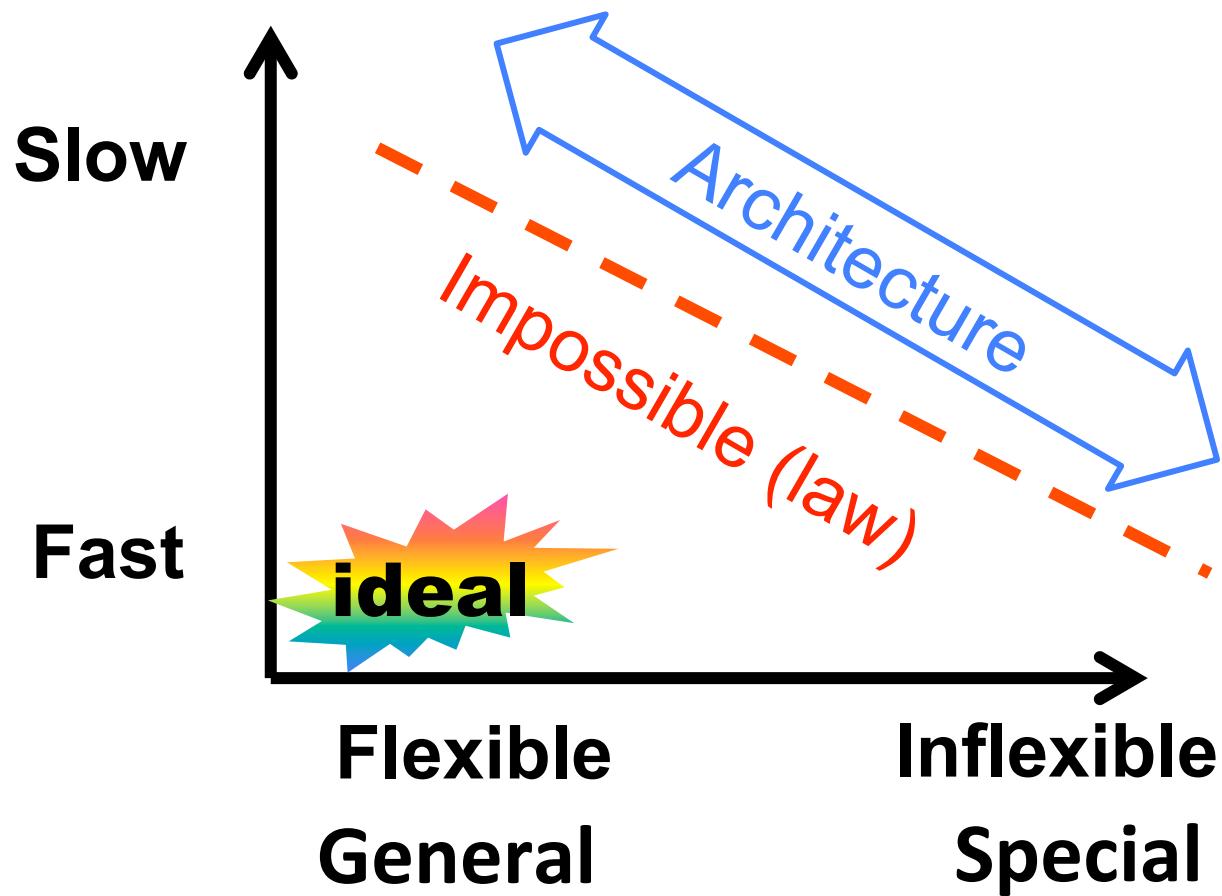
What about here?

Summary: Common Tradeoffs

- Binary machine code vs. higher-level language
- VM vs. bare metal (vs. container)
- Fast path vs. slow path
- Hardware vs. software
- Formula-1 vs. minivan
- ...
- But what are the essential features of these tradeoffs
 - What is fundamental in the tradeoff space
 - And are there “laws” governing these tradeoffs?
 - And how do they relate to RYF-complexity?

RYF tradeoffs are fundamental

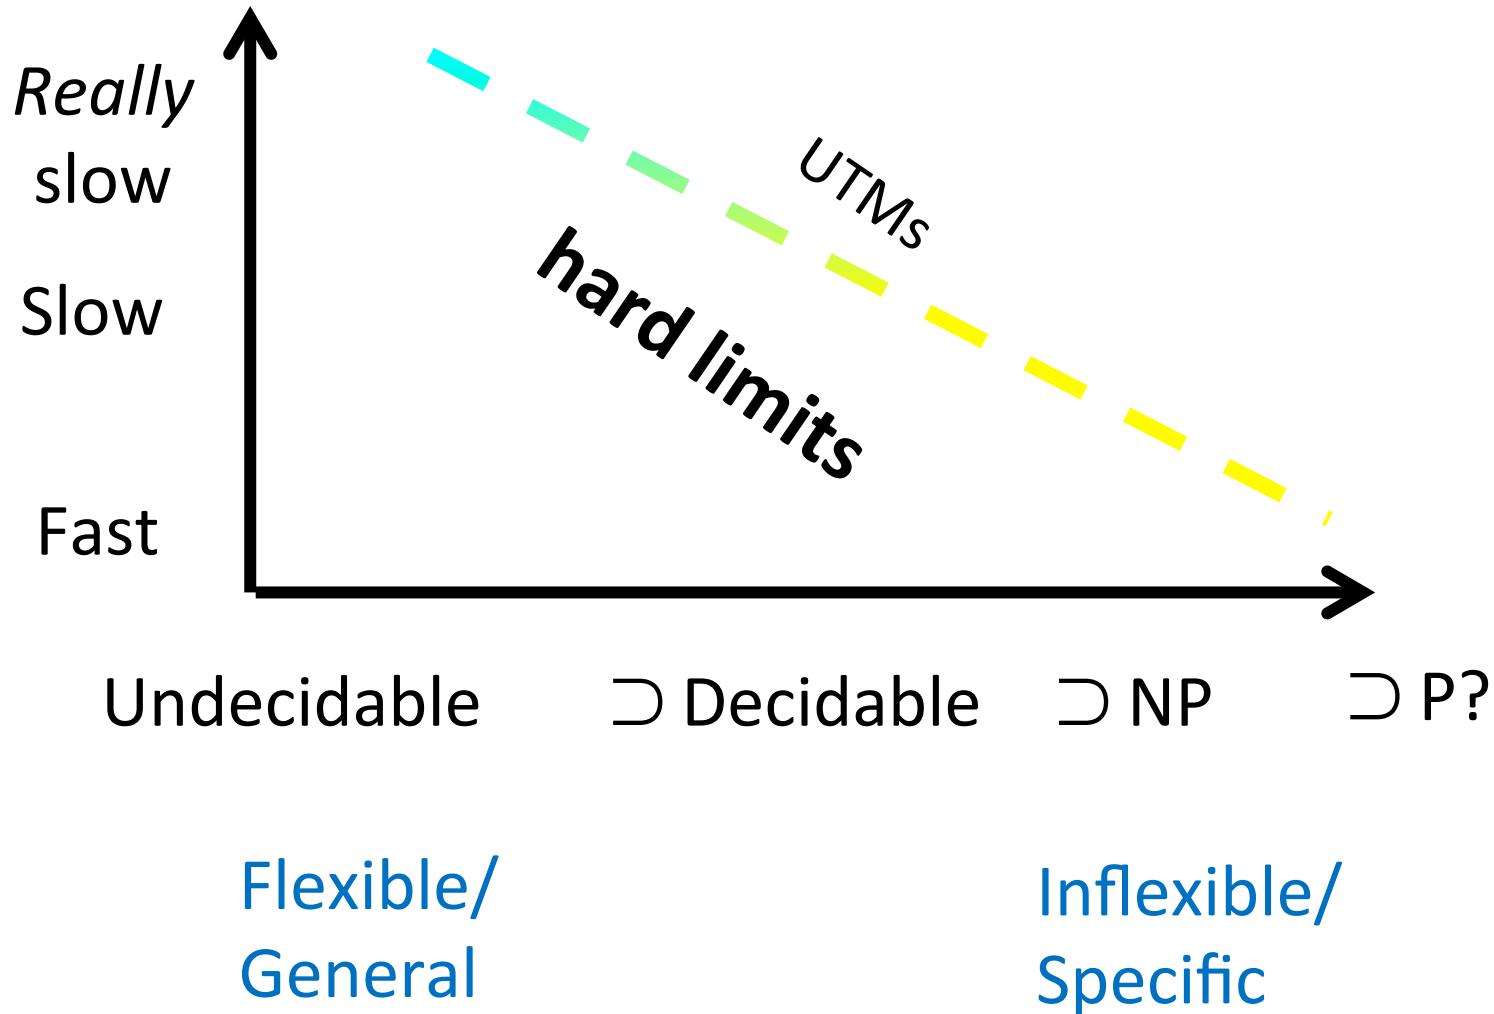
Example: Computational Complexity
Layering, Formal Systems, Hard Tradeoffs



Architecture
(constraints that deconstrain)



Drilling down a bit on the Computational Complexity Tradeoff Space

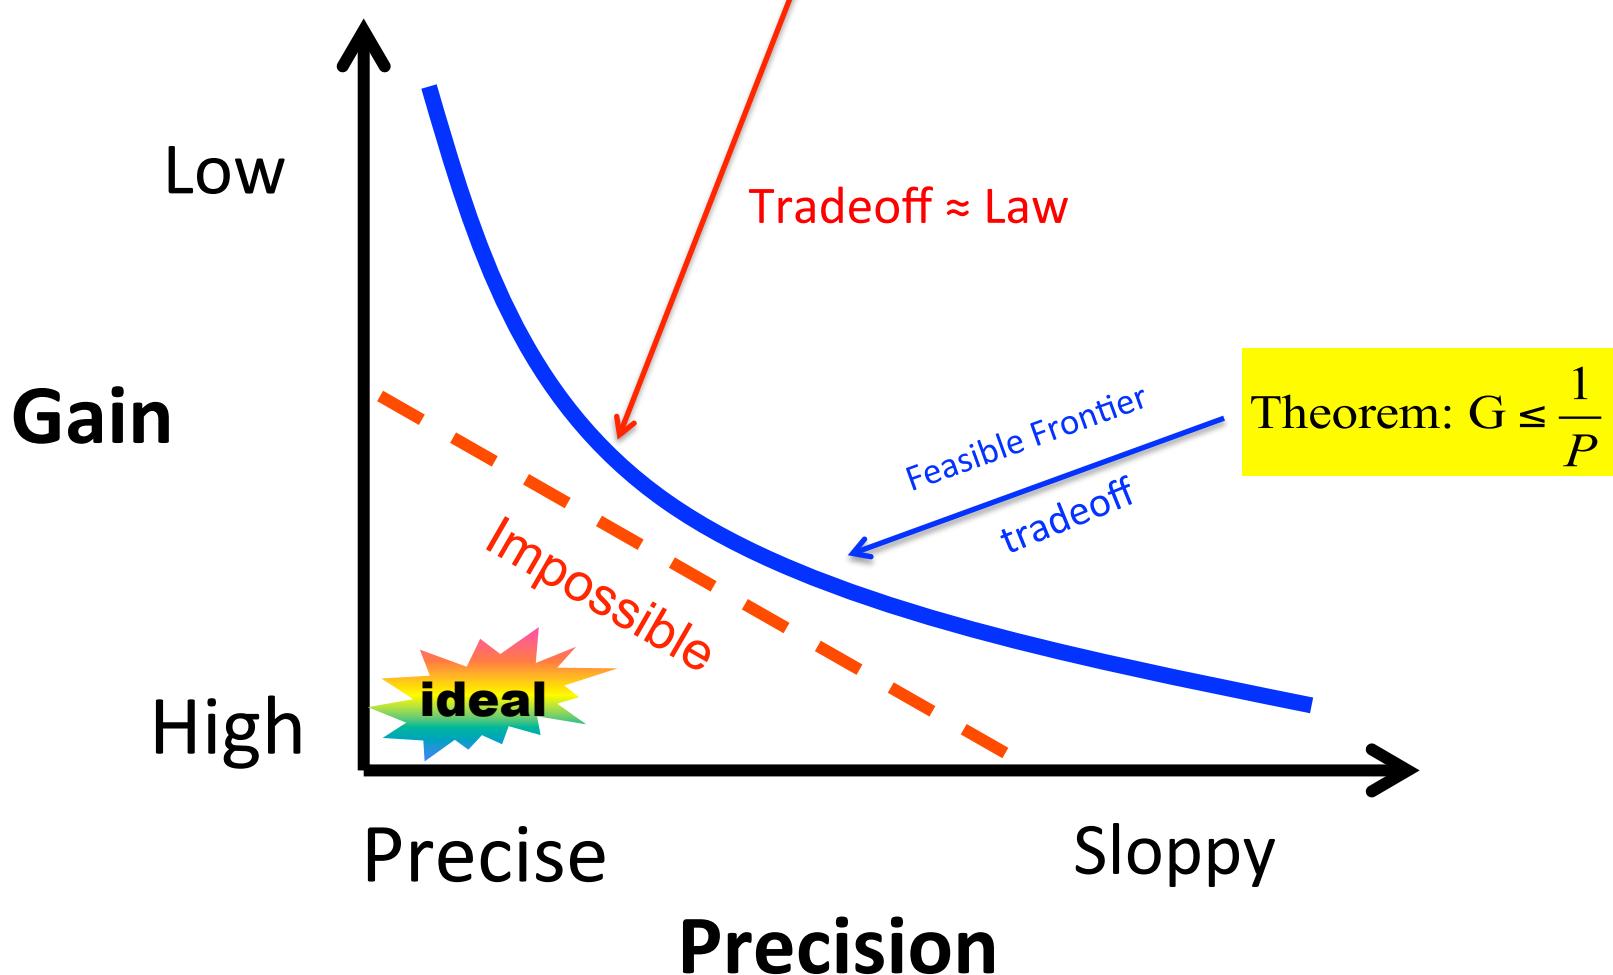


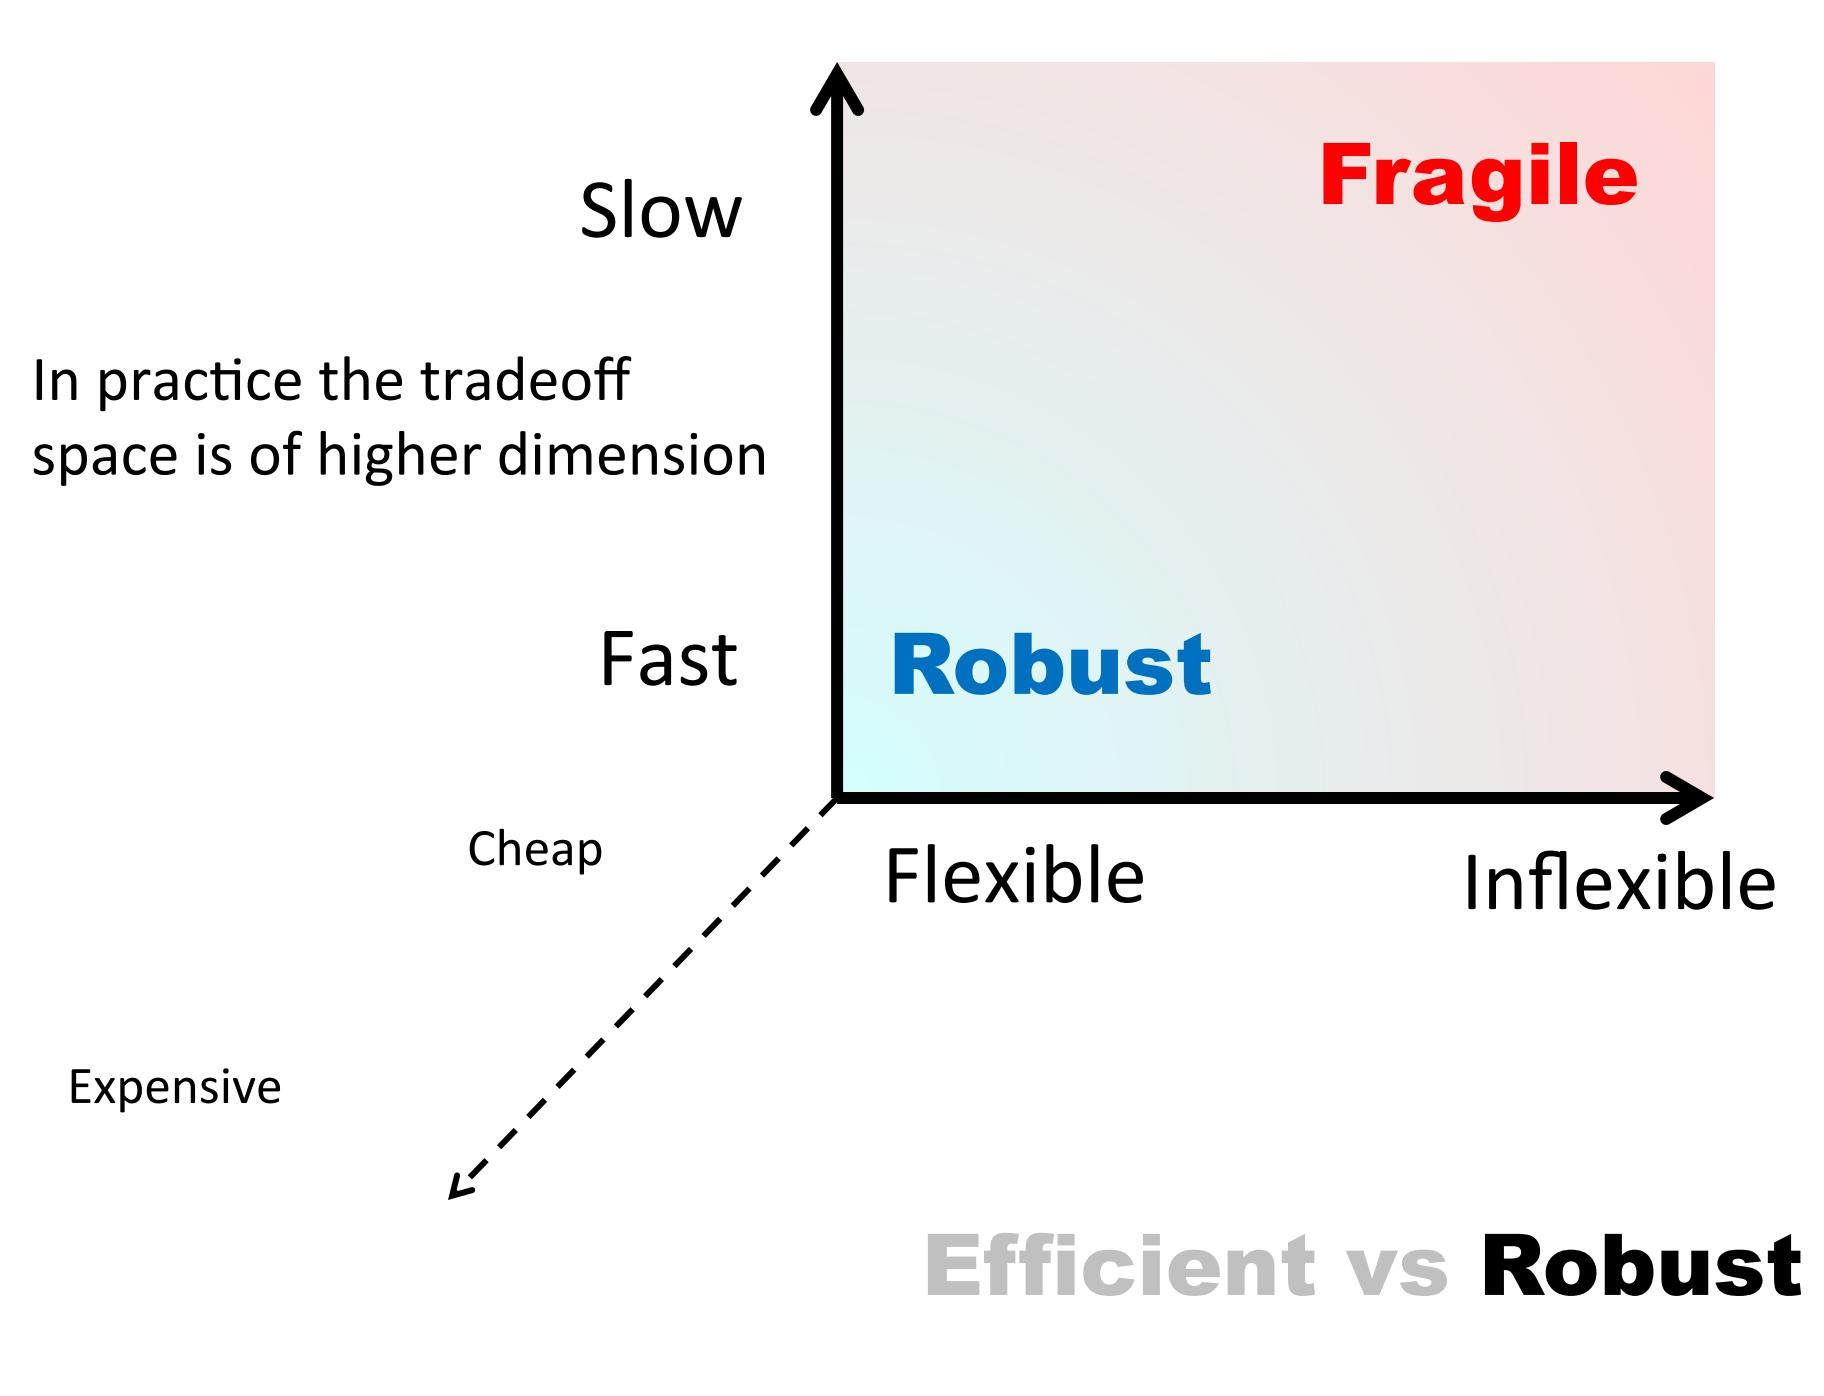
Another Example: Feedback Control Theory

(Gain/Sensitivity Tradeoff In Feedback Control)

$$\int_0^\infty \ln |S(i\omega)| d\omega = \int_0^\infty \ln \left| \frac{1}{1 + L(i\omega)} \right| d\omega = \pi \sum \operatorname{Re}(p_k) - \frac{\pi}{2} \lim_{s \rightarrow \infty} sL(s)$$

Bode Sensitivity Integral





Agenda

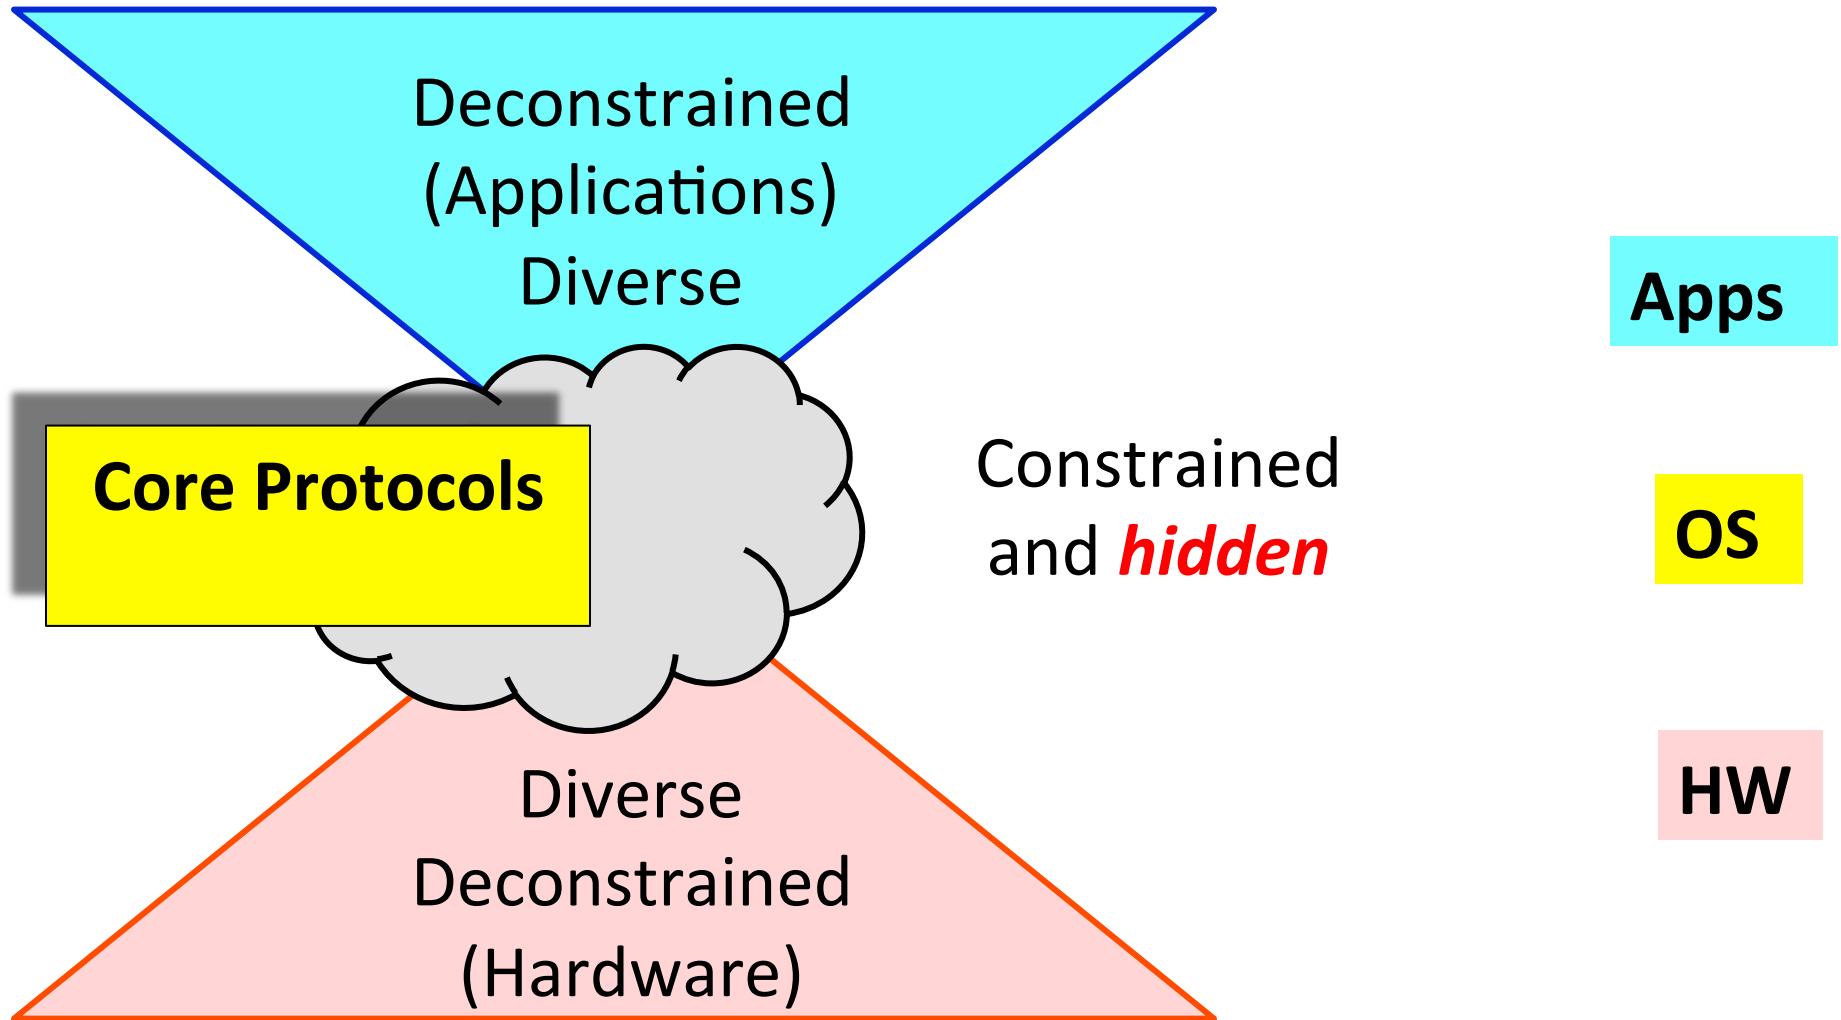
- ~~Goals for this Talk~~
- ~~What is Complexity and Why is it Hidden~~
 - ~~Robustness, Fragility, and Complexity~~
- ~~Tradeoffs?~~
- The Architecture of Complex Systems
 - Universal Architectural Principles
- The Evolution of DevOps
- A Few Conclusions and Q&A if we have time

The Architecture of Complex Systems

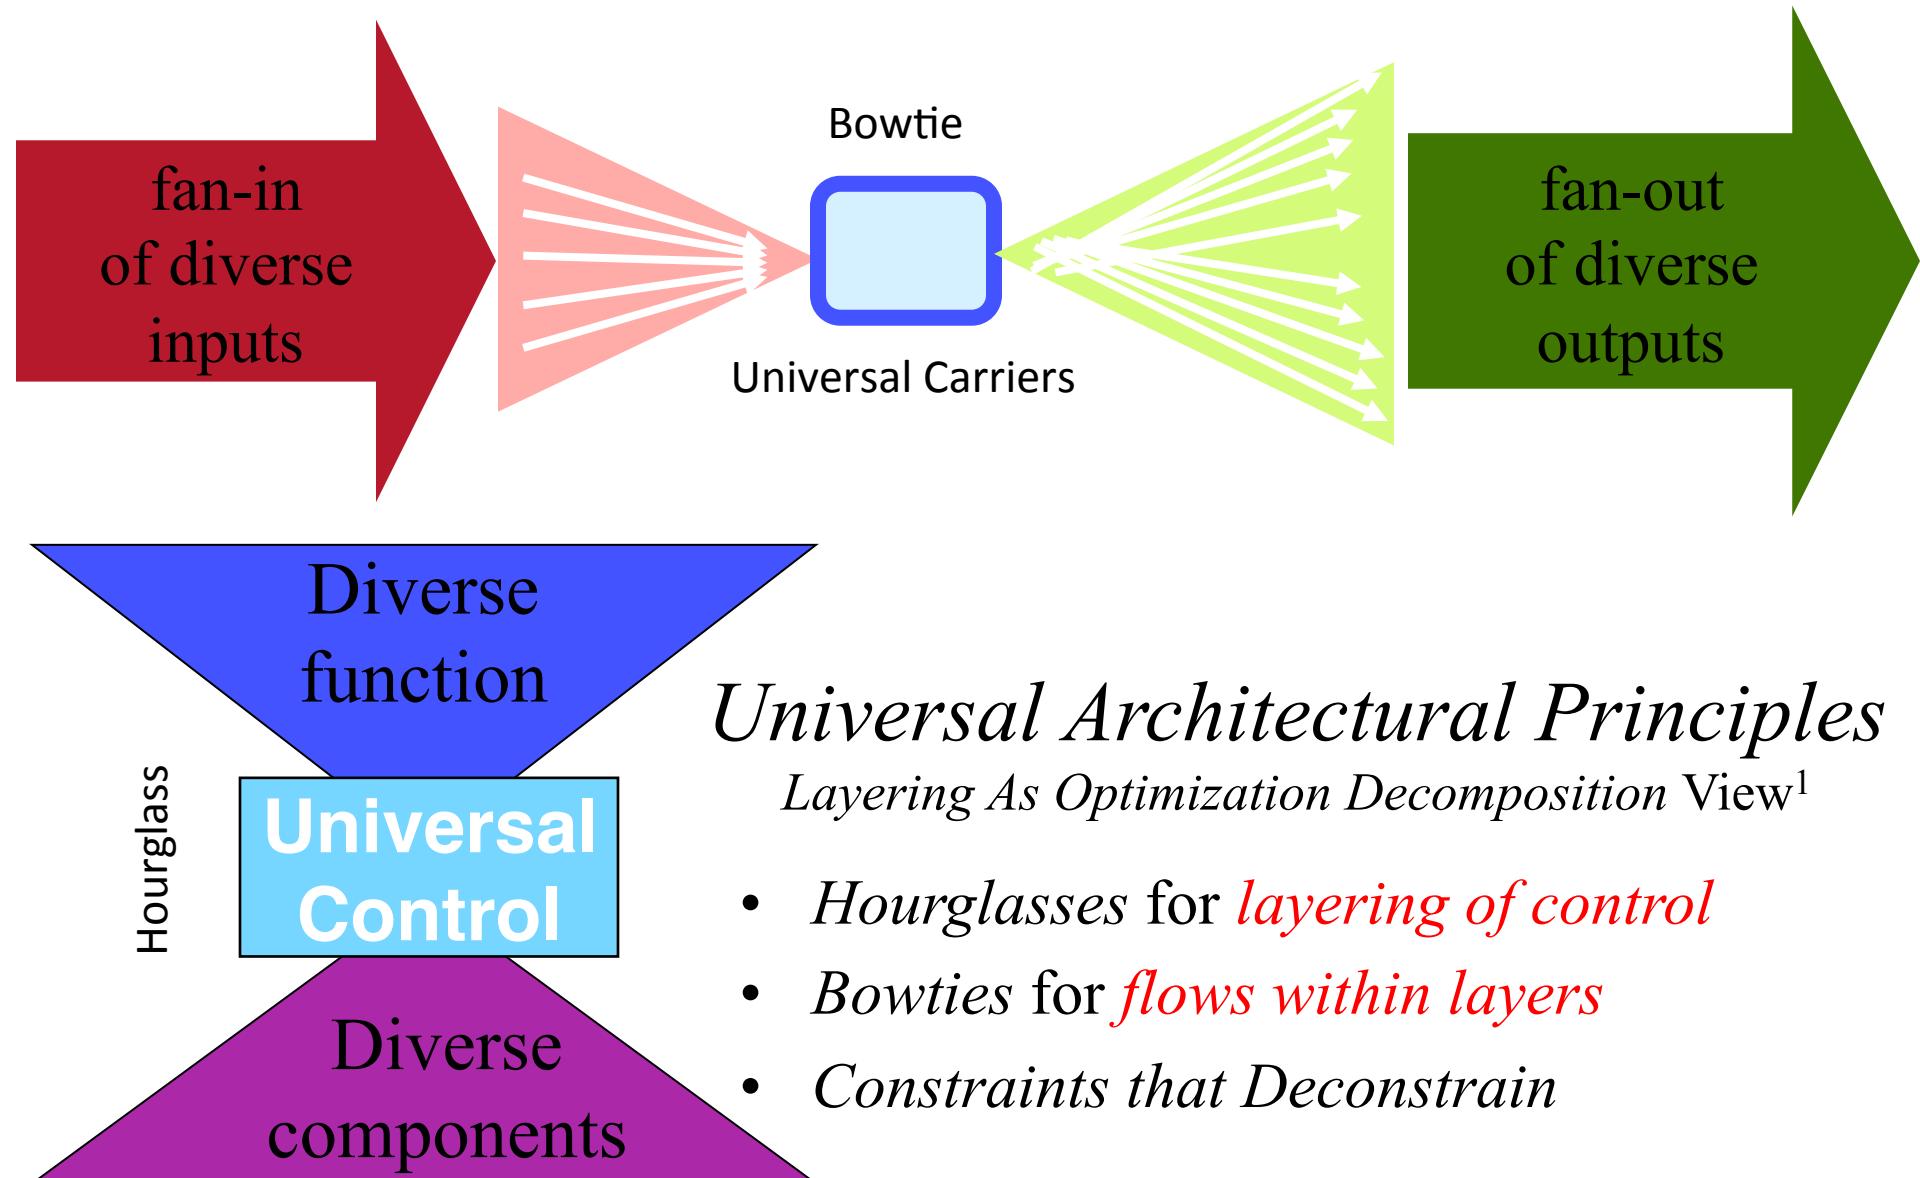
- What we have learned is that there are *universal architectural building blocks* found in systems that scale and are evolvable. These include
- **Architecture/Layering**
- **Laws, constraints, tradeoffs**
- Protocol Based Architectures
- Massively distributed with *robust* control loops
- Consequences
 - Hidden RYF Complexity
 - Hijacking, parasitism, predation

So What is the Basic Layered Architecture?

(hint: layering is *the* fundamental architectural feature)

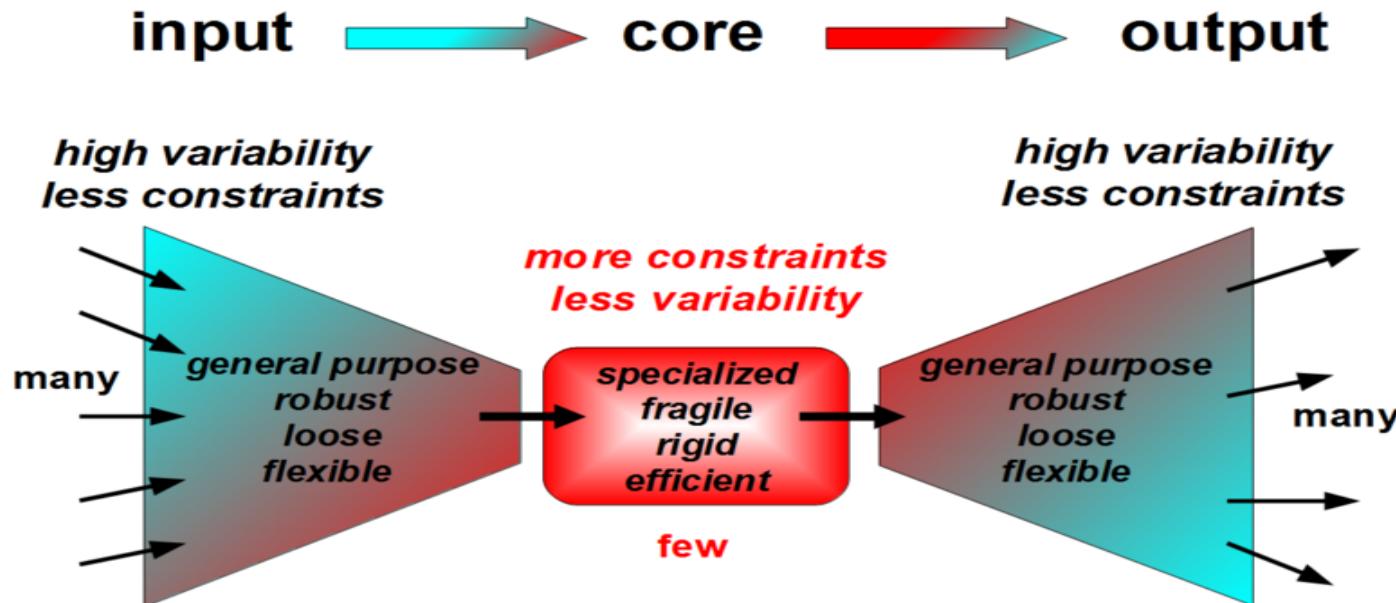


Bowties, Hourglasses and Layering



Bowties 101

Constraints that Deconstrain Schematic of a “Layer”

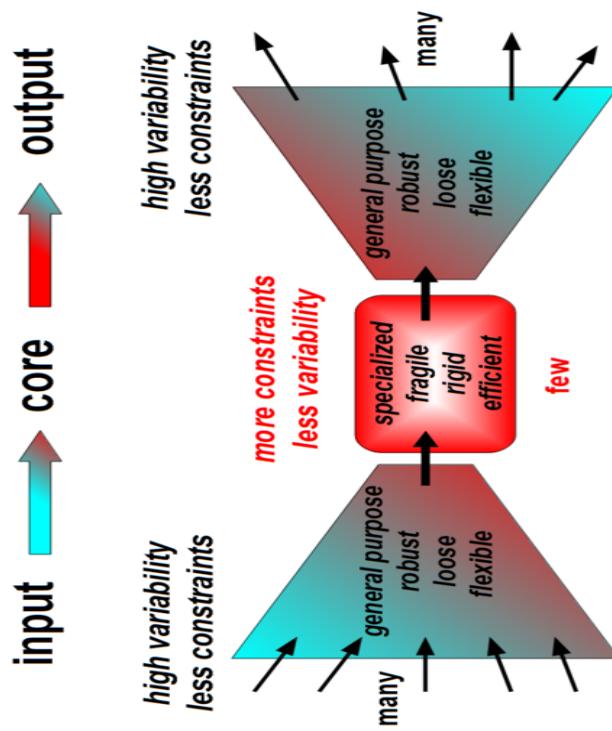


For example, the reactions and metabolites of core metabolism, e.g., Adenosine Triphosphate (ATP) metabolism, Krebs/Citric Acid Cycle, ... form a “metabolic knot”. That is, ATP is a **Universal Carrier** for cellular energy.

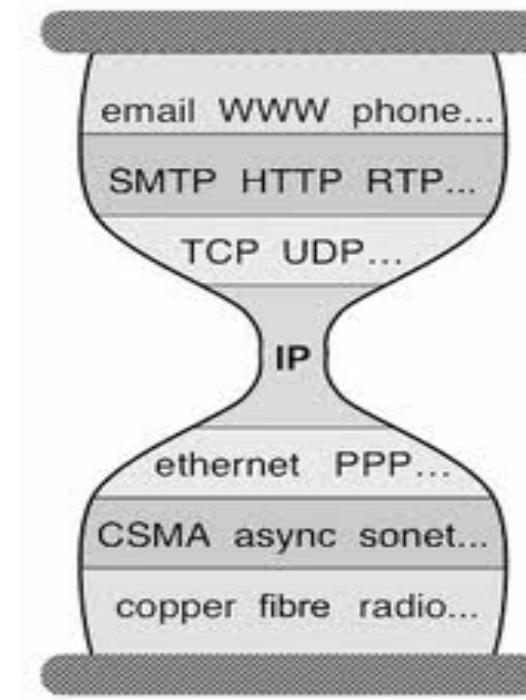
1. Processes L-1 information and/or raw material flows into a “standardized” format (the L+1 abstraction)
2. Provides plug-and-play modularity for the layer above
3. Provides robustness but at the same time fragile to attacks against/using the standardized interface

But Wait a Second

Anything Look Familiar?



Bowtie Architecture

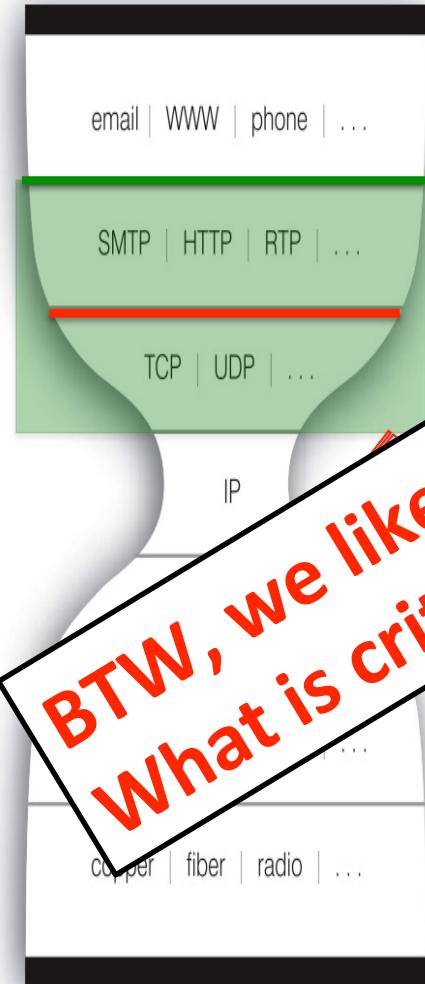


Hourglass Architecture

Comes down to whether you see layering as horizontal or vertical

The Nested Bowtie/Hourglass Architecture of the Internet

Layering of Control



HTTP Bowtie

Input: Ports, Datagrams, Connections

Output (abstraction): REST

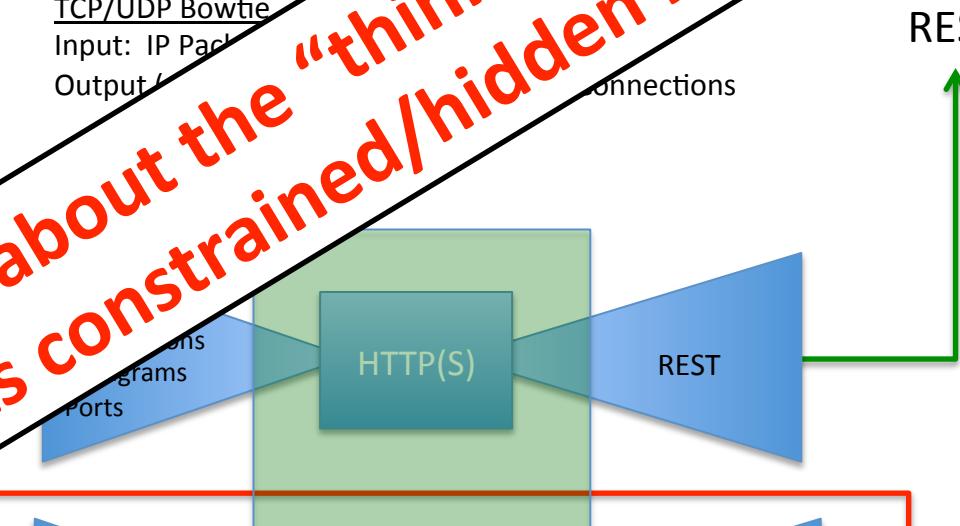
TCP/UDP Bowtie

Input: IP Packets, Ports, Connections

Output (abstraction): REST

Layering of Control/Abstractions

**BTW, we like to talk about the “thinness” of the waist
What is critical is its constrained/hidden nature**



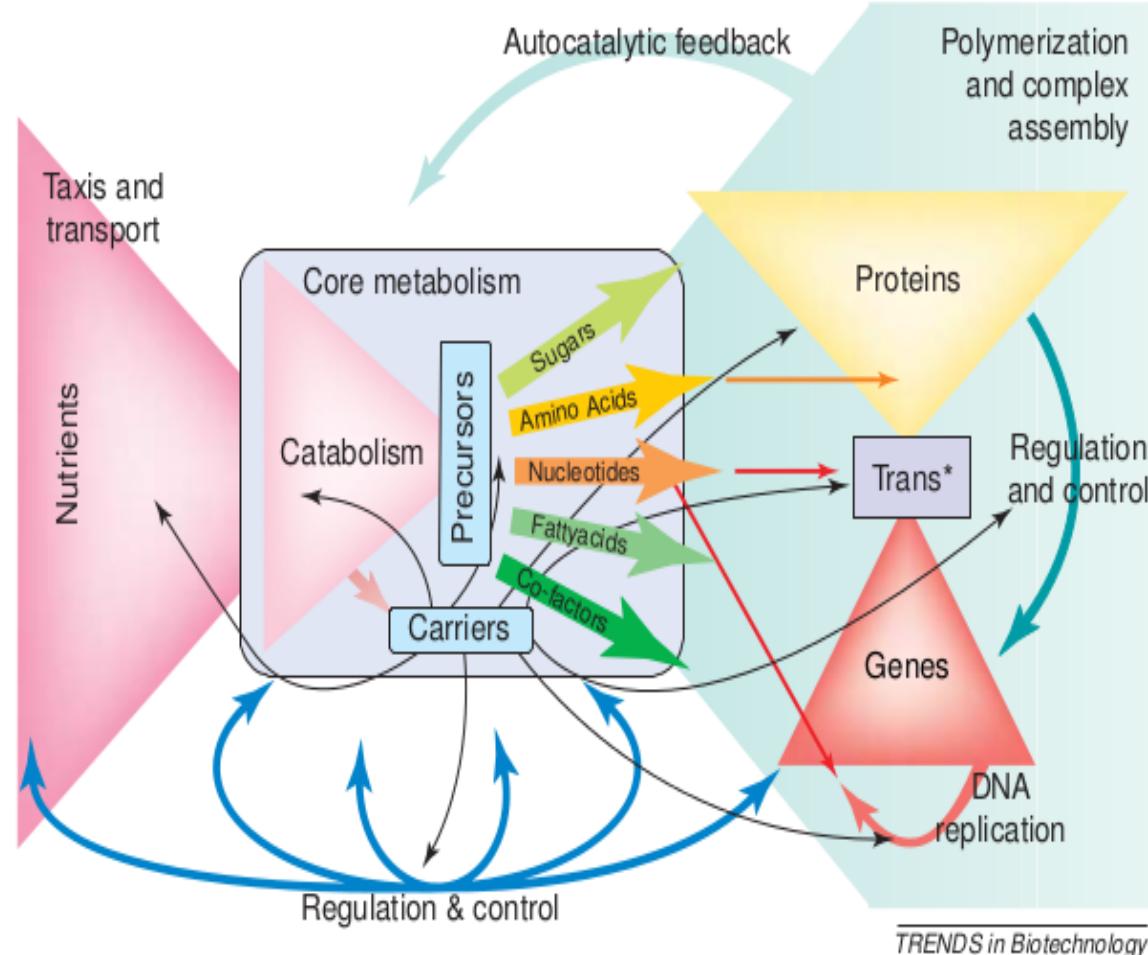
$$\begin{aligned} & \text{maximize} && \sum_s U_s(x_s) \\ & \text{subject to} && Rx \leq c. \end{aligned}$$

Flows within Layers

Reverse/forward engineering Network Utility Maximum (NUM) problem
(aka Monotropic Programming)

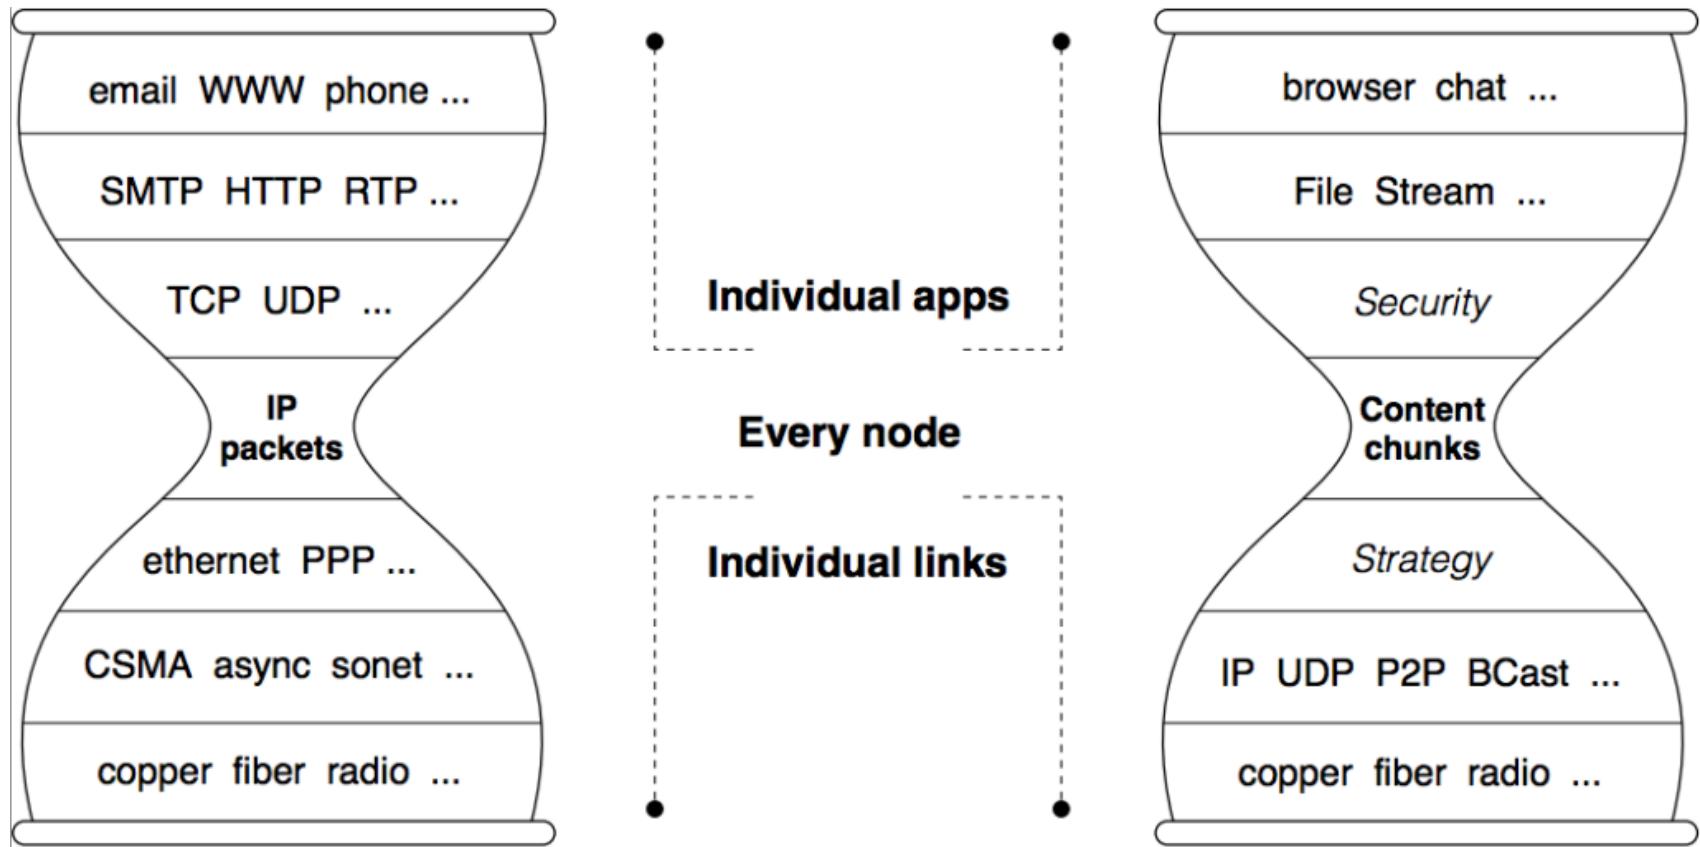
In Practice Things are More Complicated

The Nested Bowtie/Hourglass Architecture of Metabolism

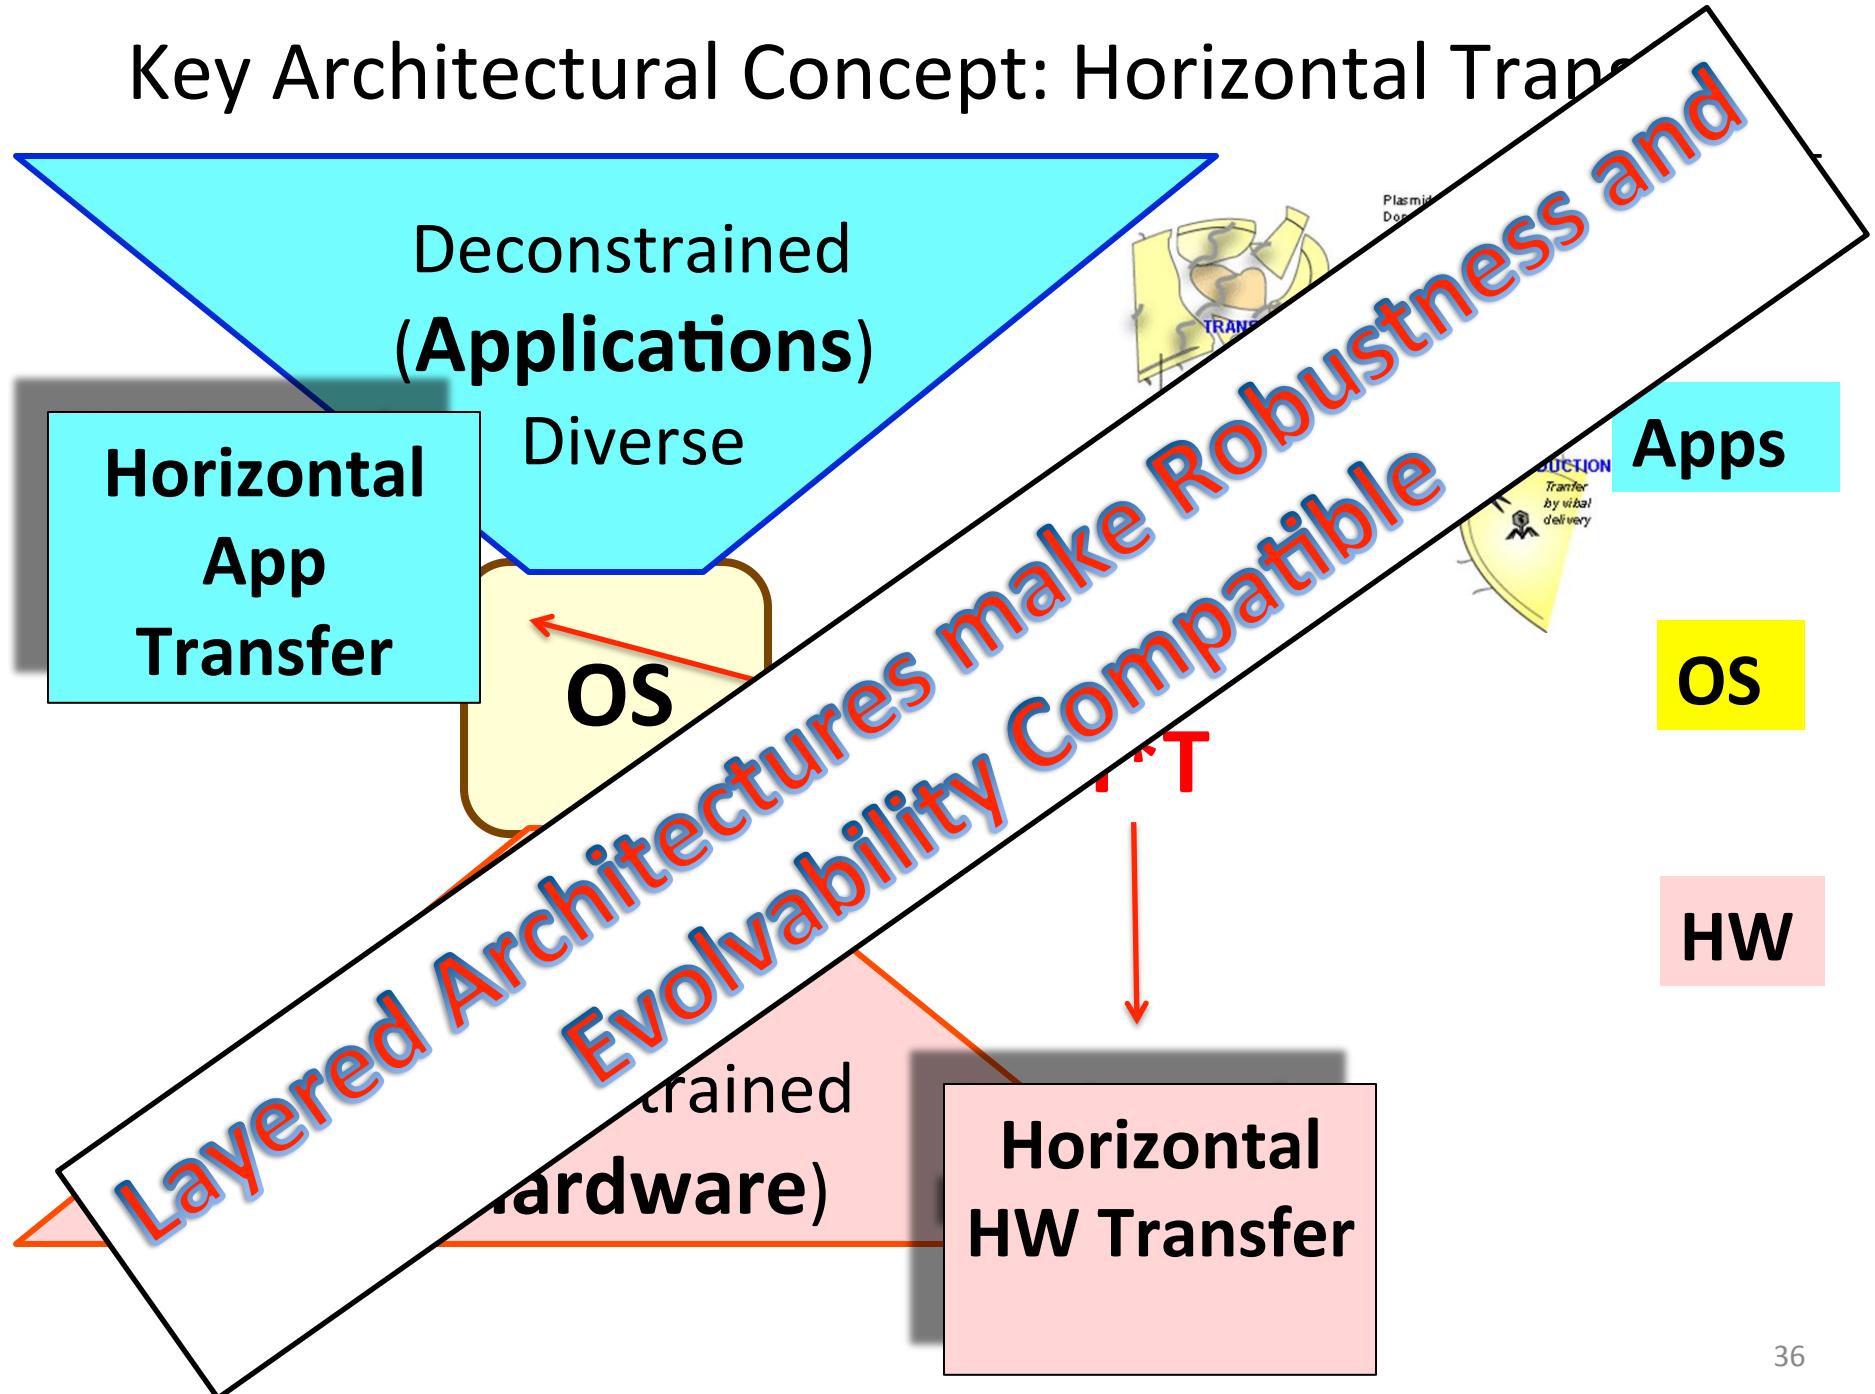


TRENDS in Biotechnology

Named Data Networking Hourglass



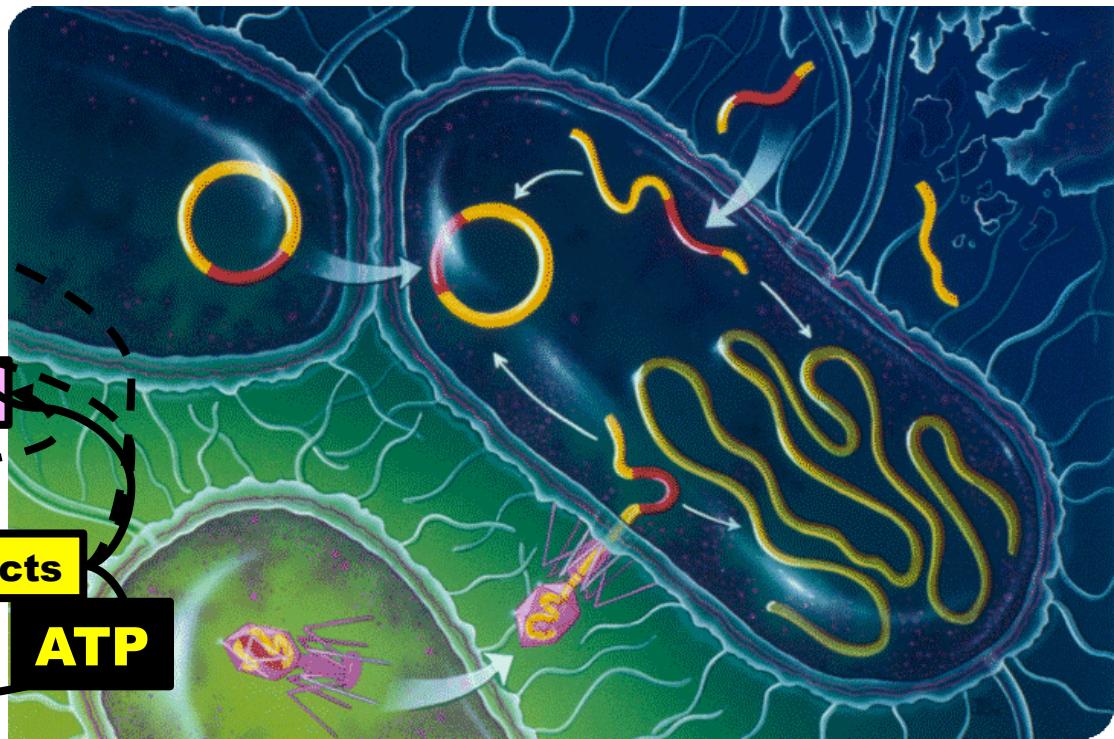
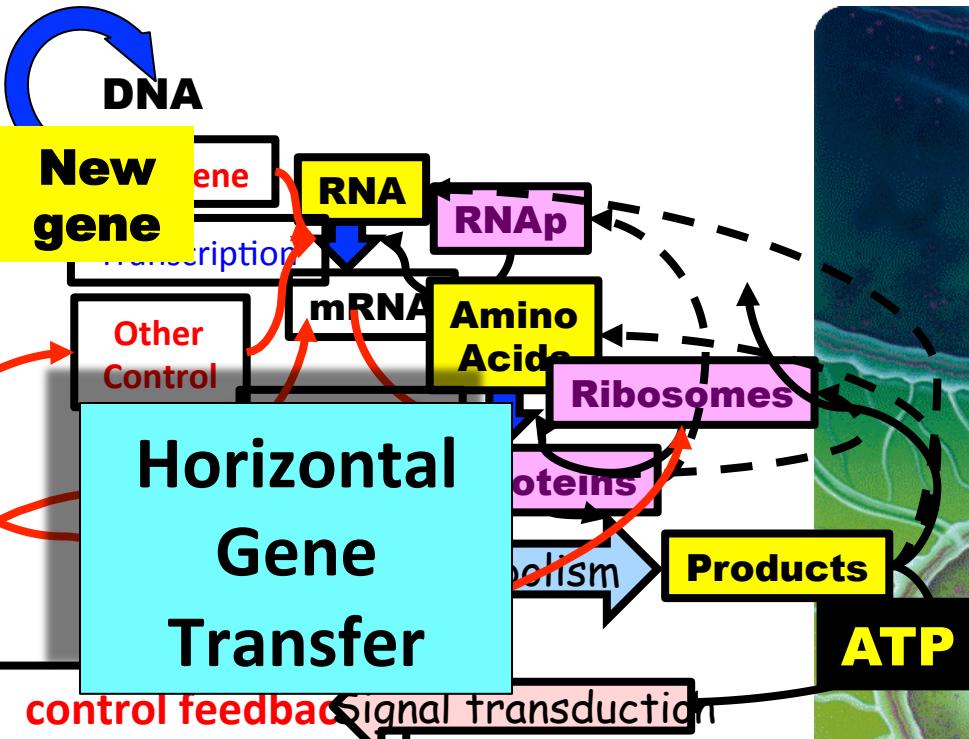
Key Architectural Concept: Horizontal Transfer



Drilling Down: HGT and antibiotic resistance

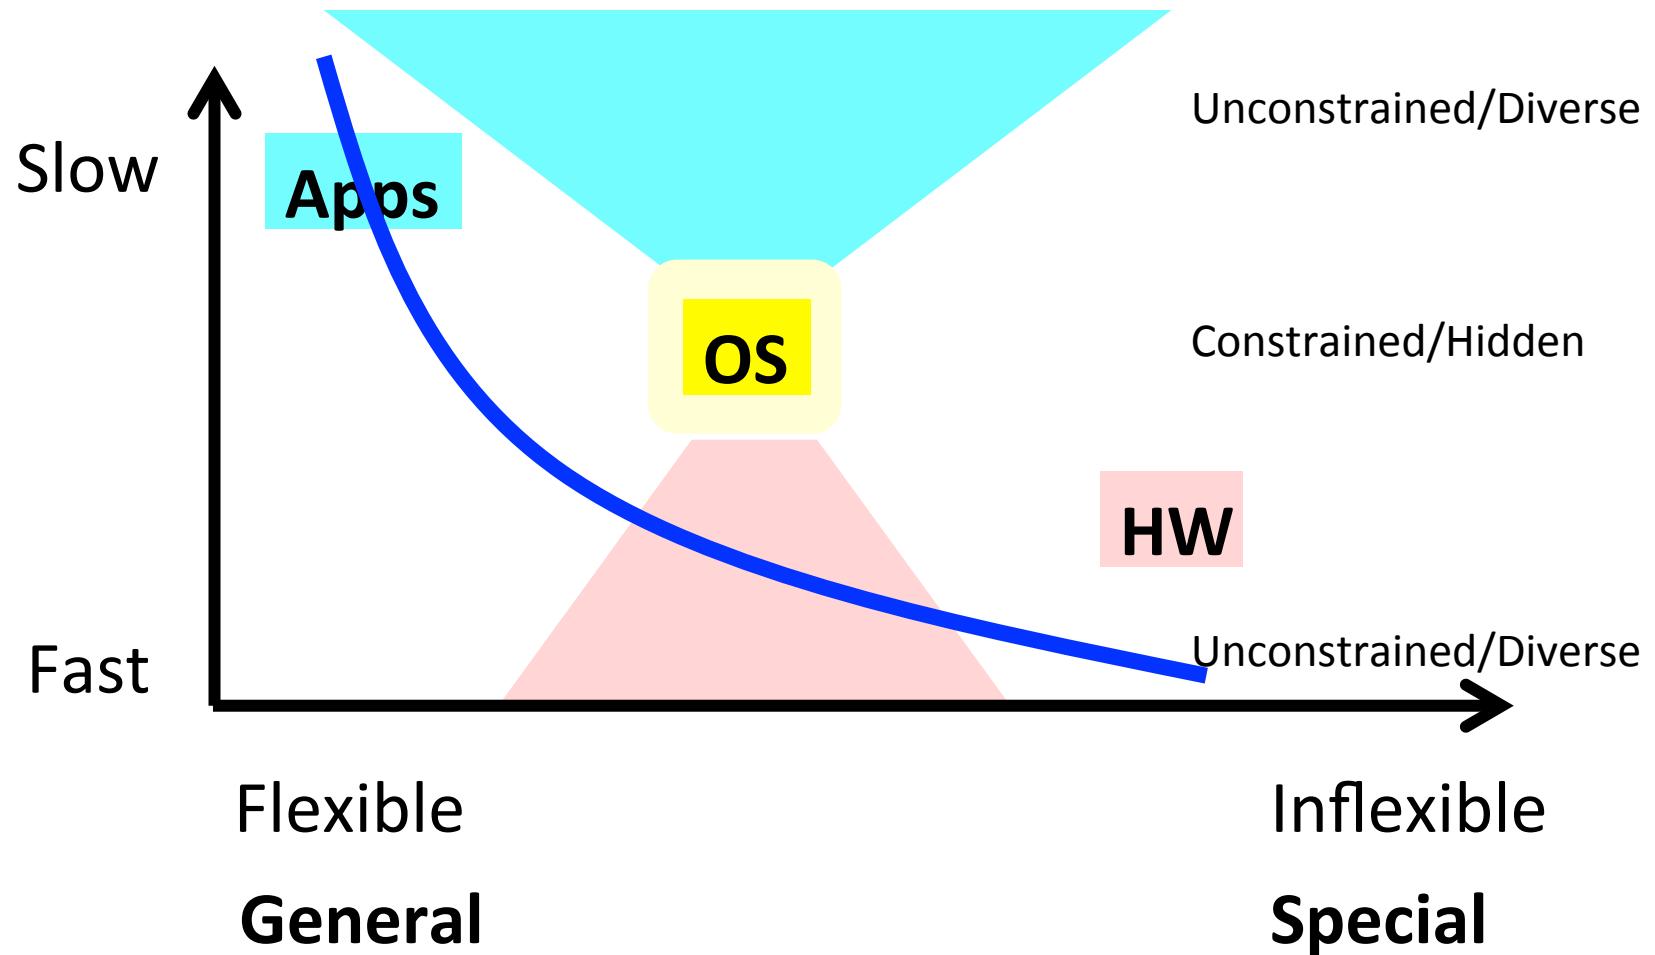
Sequence ~100 E Coli (*not* chosen randomly)

- ~ 4K genes per cell
- ~20K *different* genes in total (pan genome)
- ~ 1K universally shared genes
- ~ 300 essential (minimal) genes

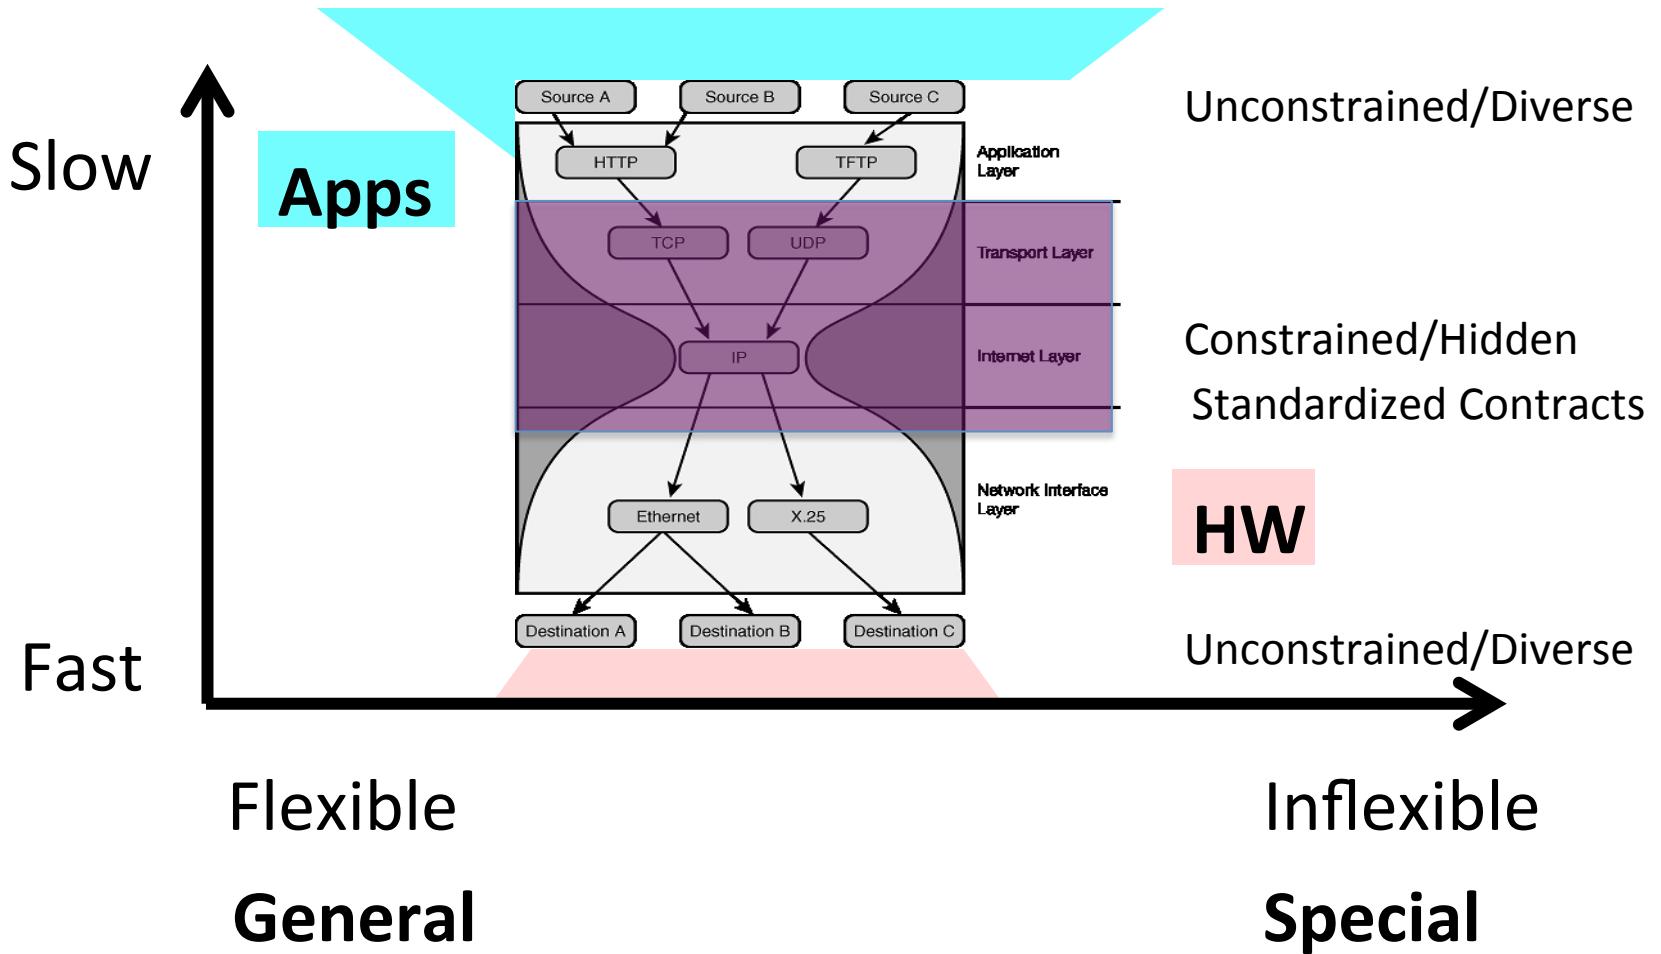


Putting it all Together

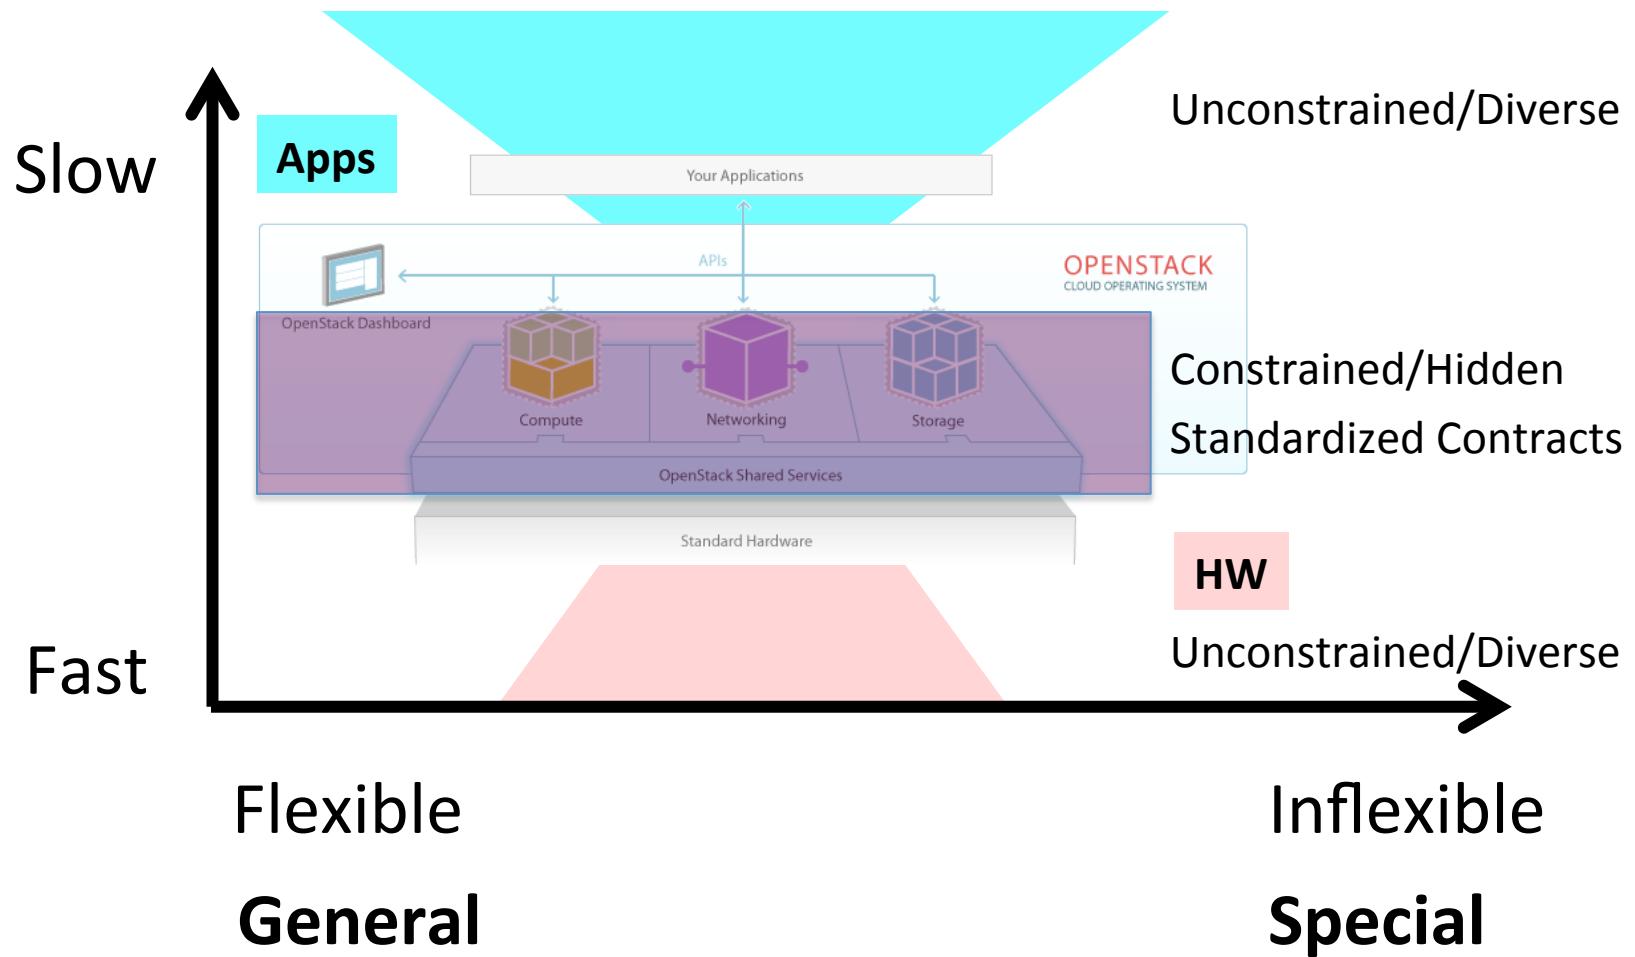
Architecture, Layering, and Tradeoffs



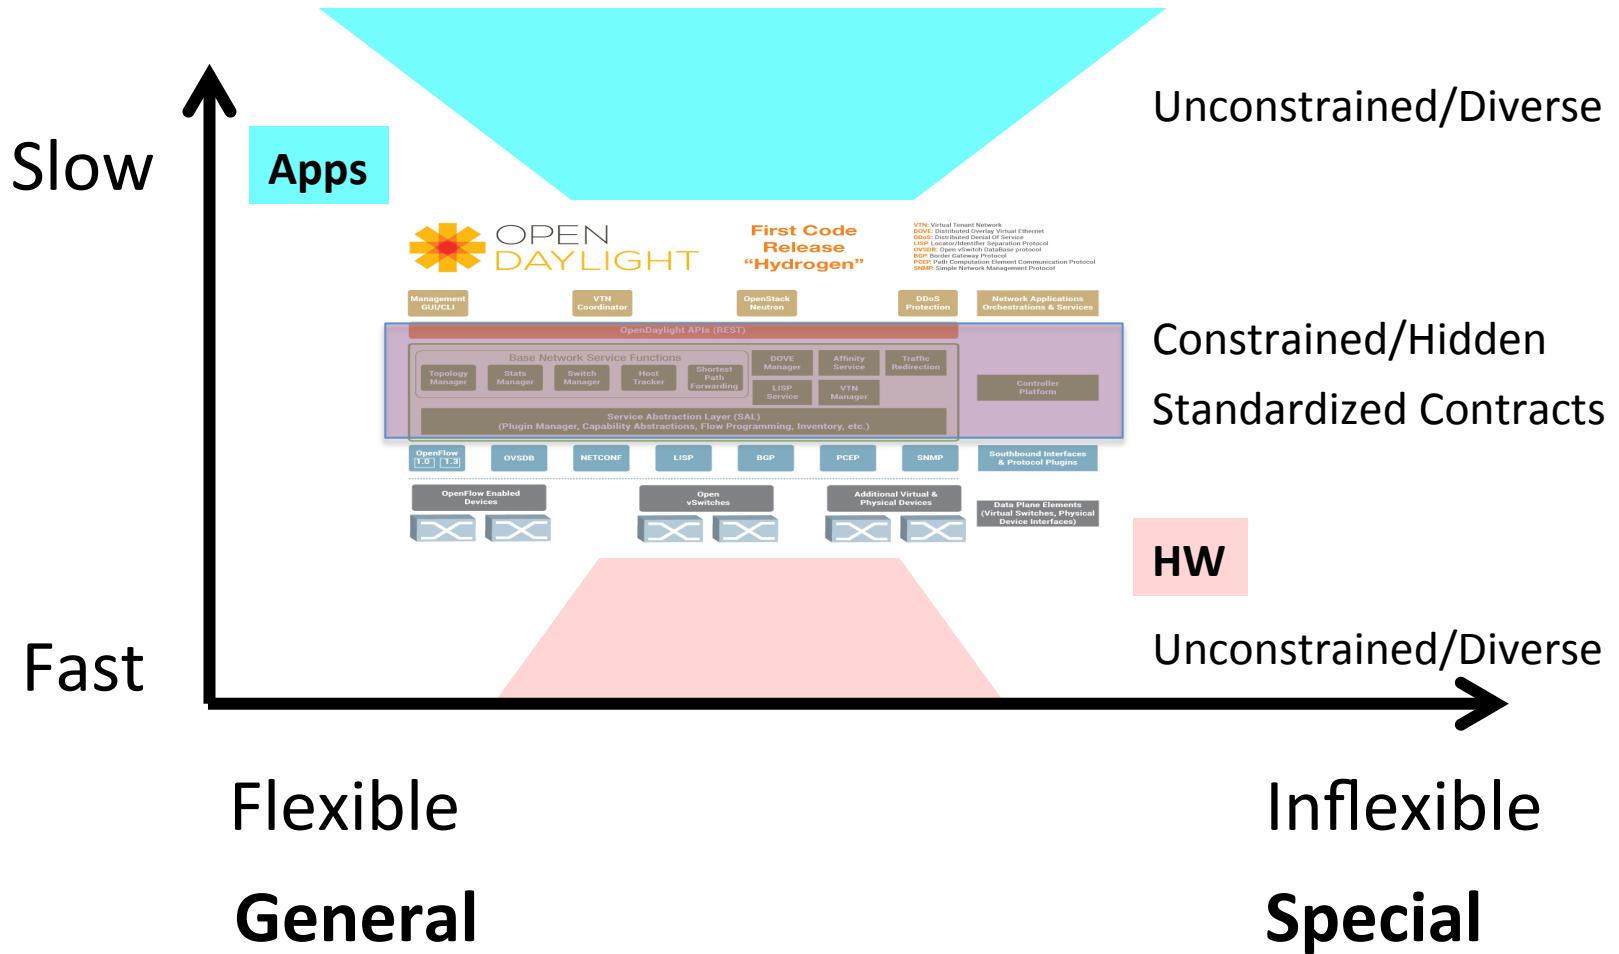
Example: Layered Network Architectures (Internet)



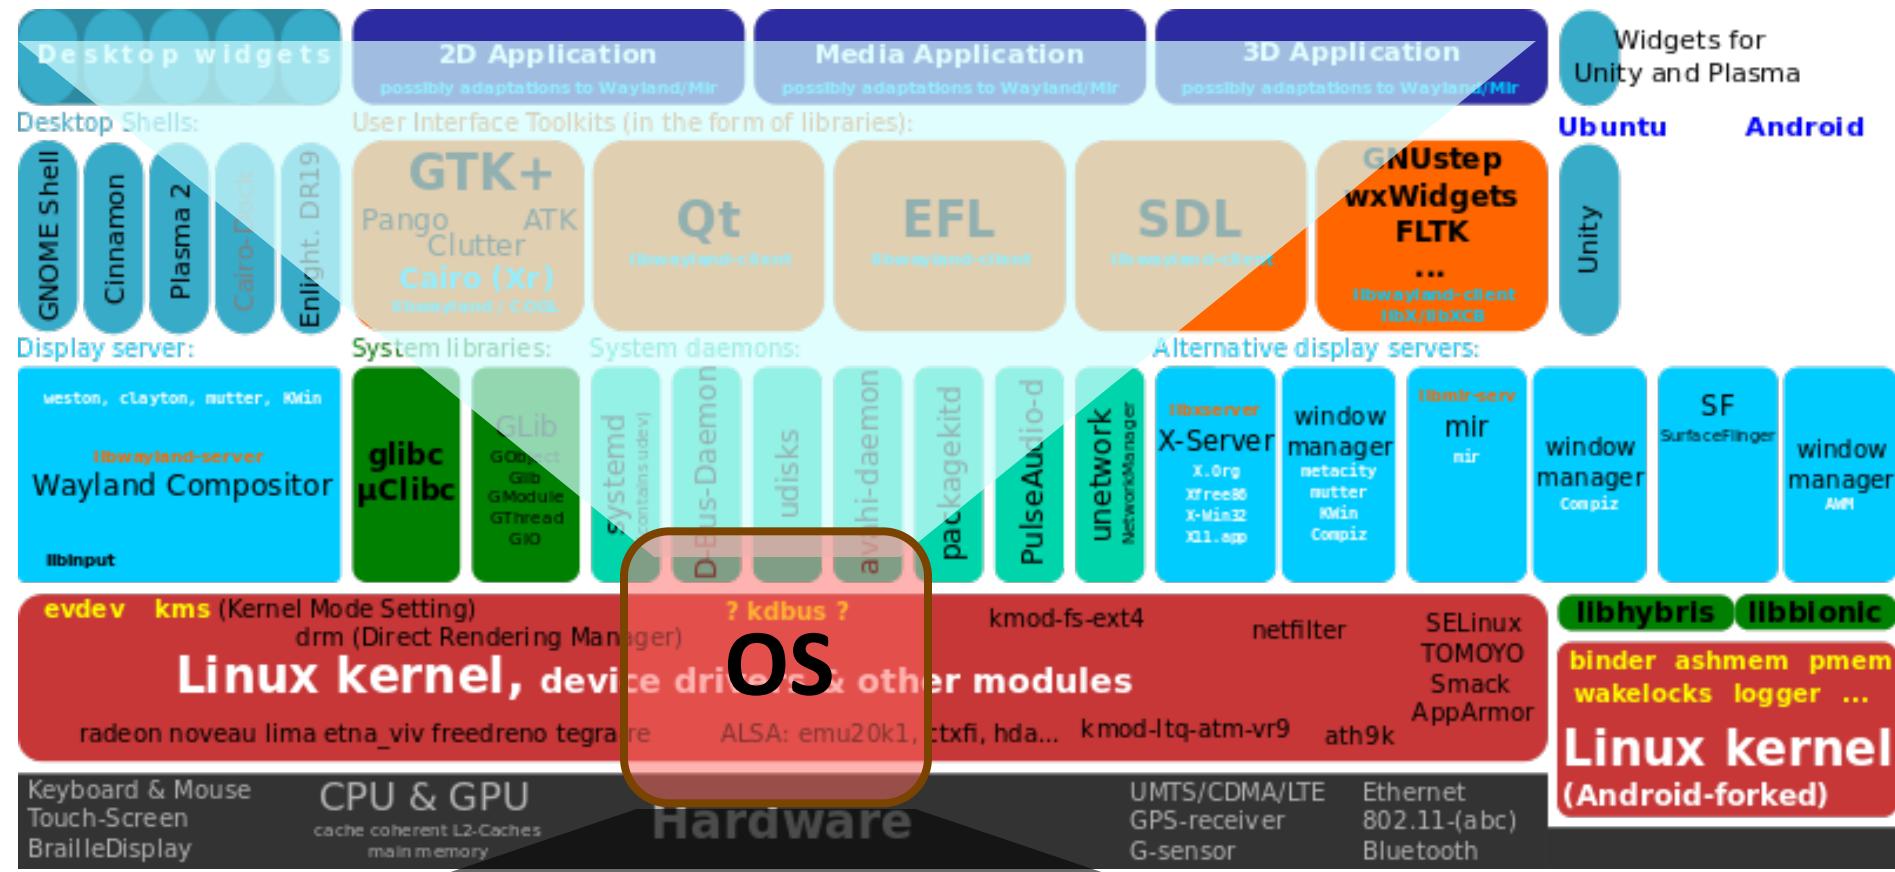
Example: OpenStack



Example: SDN



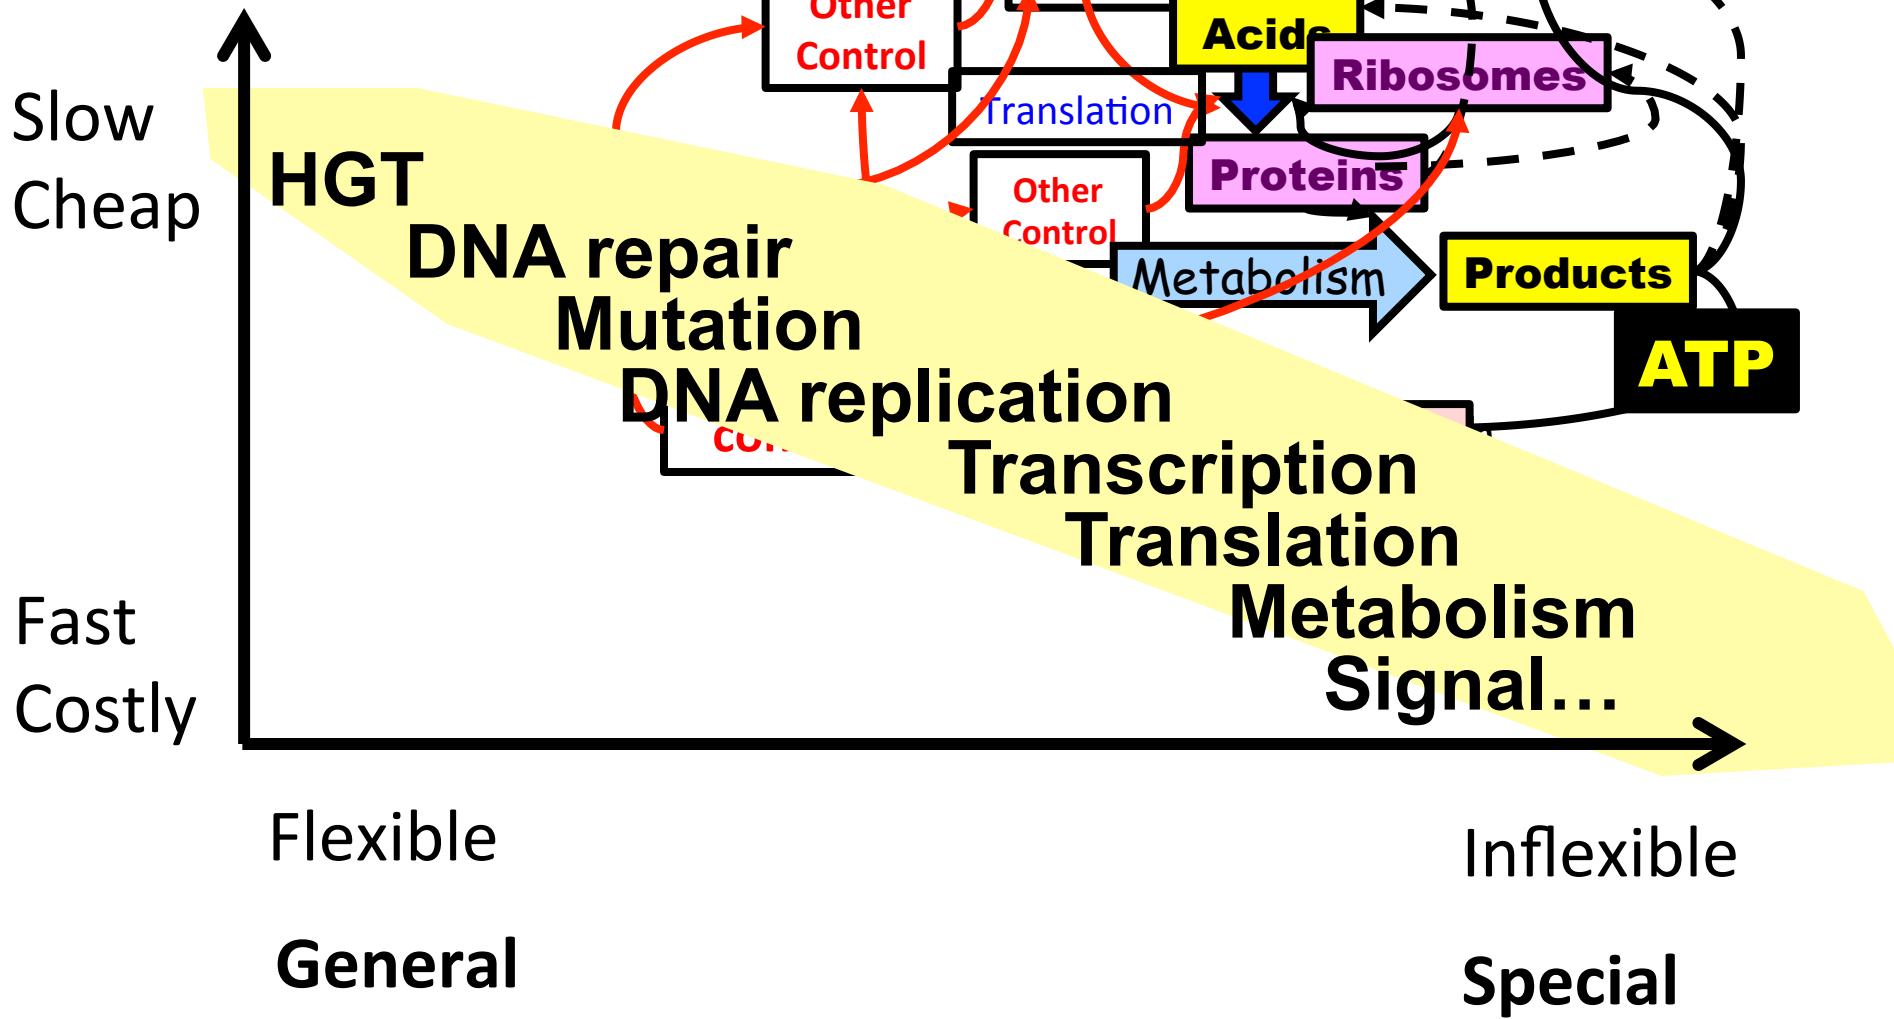
Linux Kernel?



So I've been talking a bit about Biology

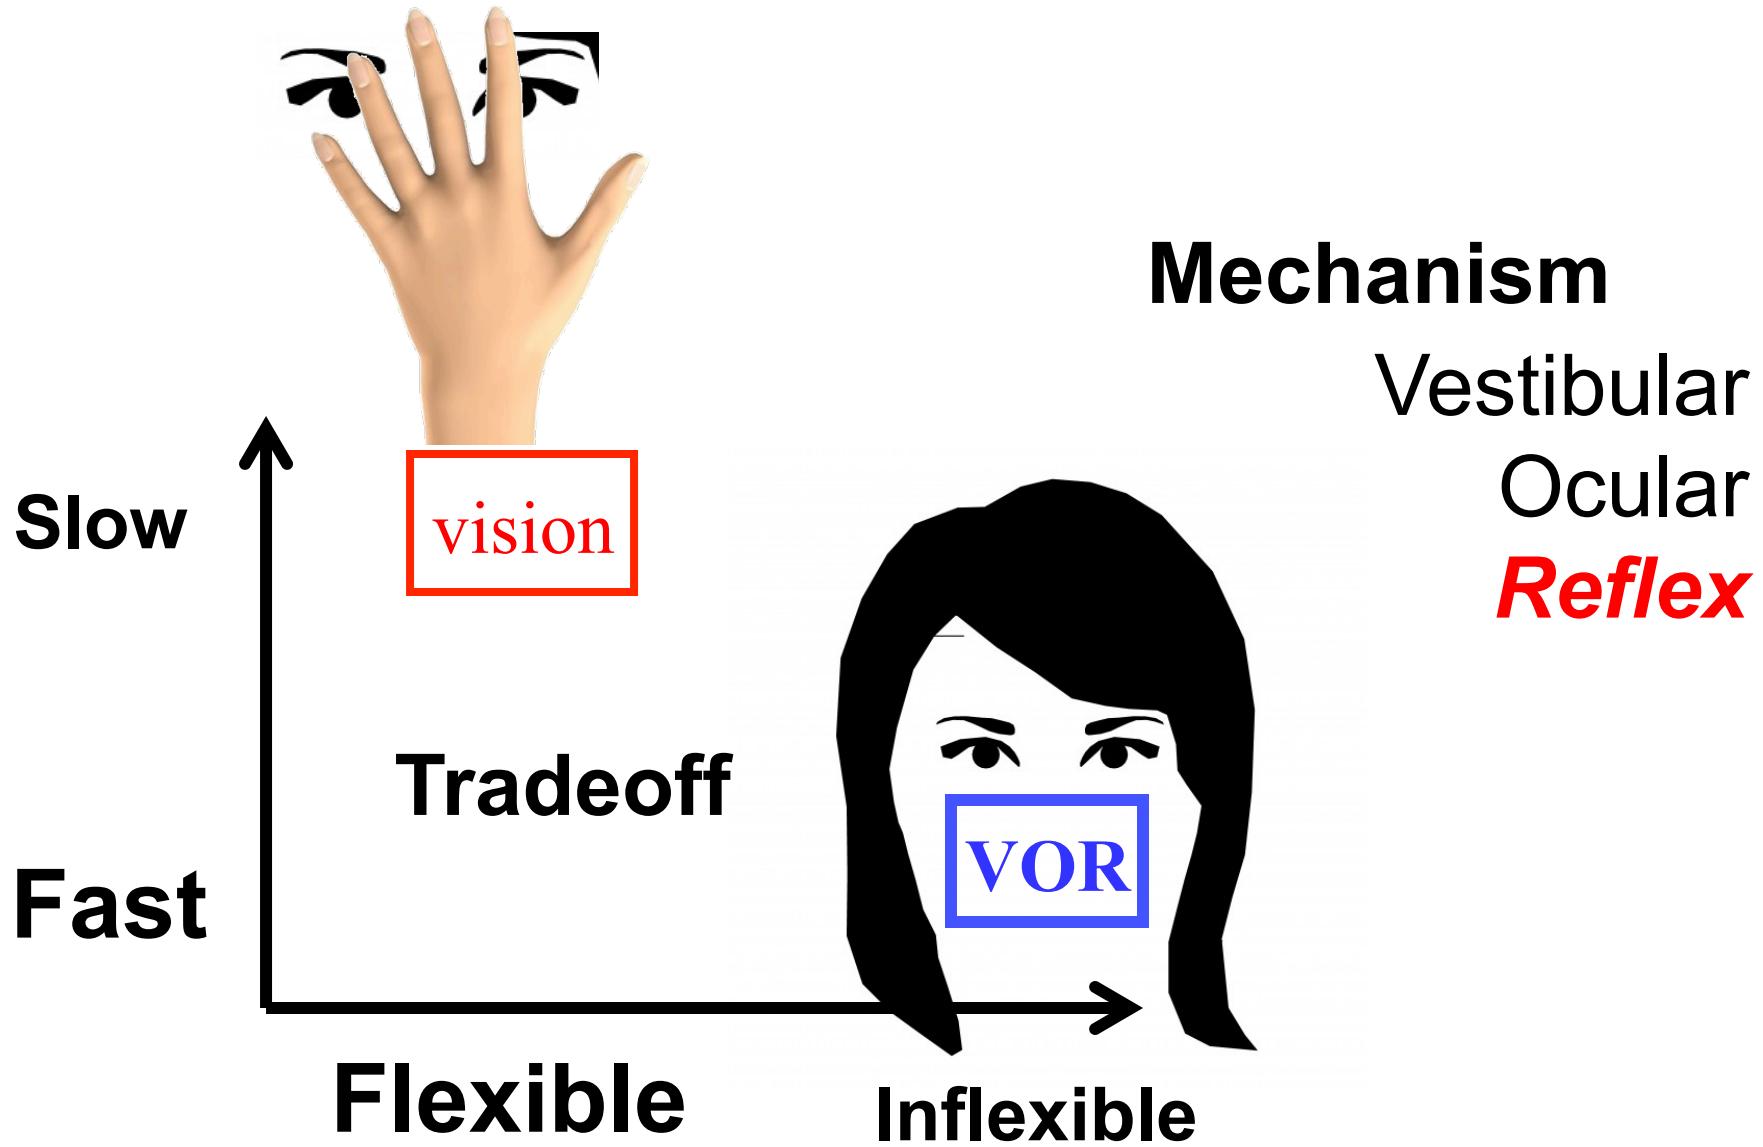
- Are biological systems architecturally layered in the same way?
 - Basically diverse apps and h/w, hidden kernel
 - Constraints that deconstrain
- If so, what kinds of systems can we analyze this way?

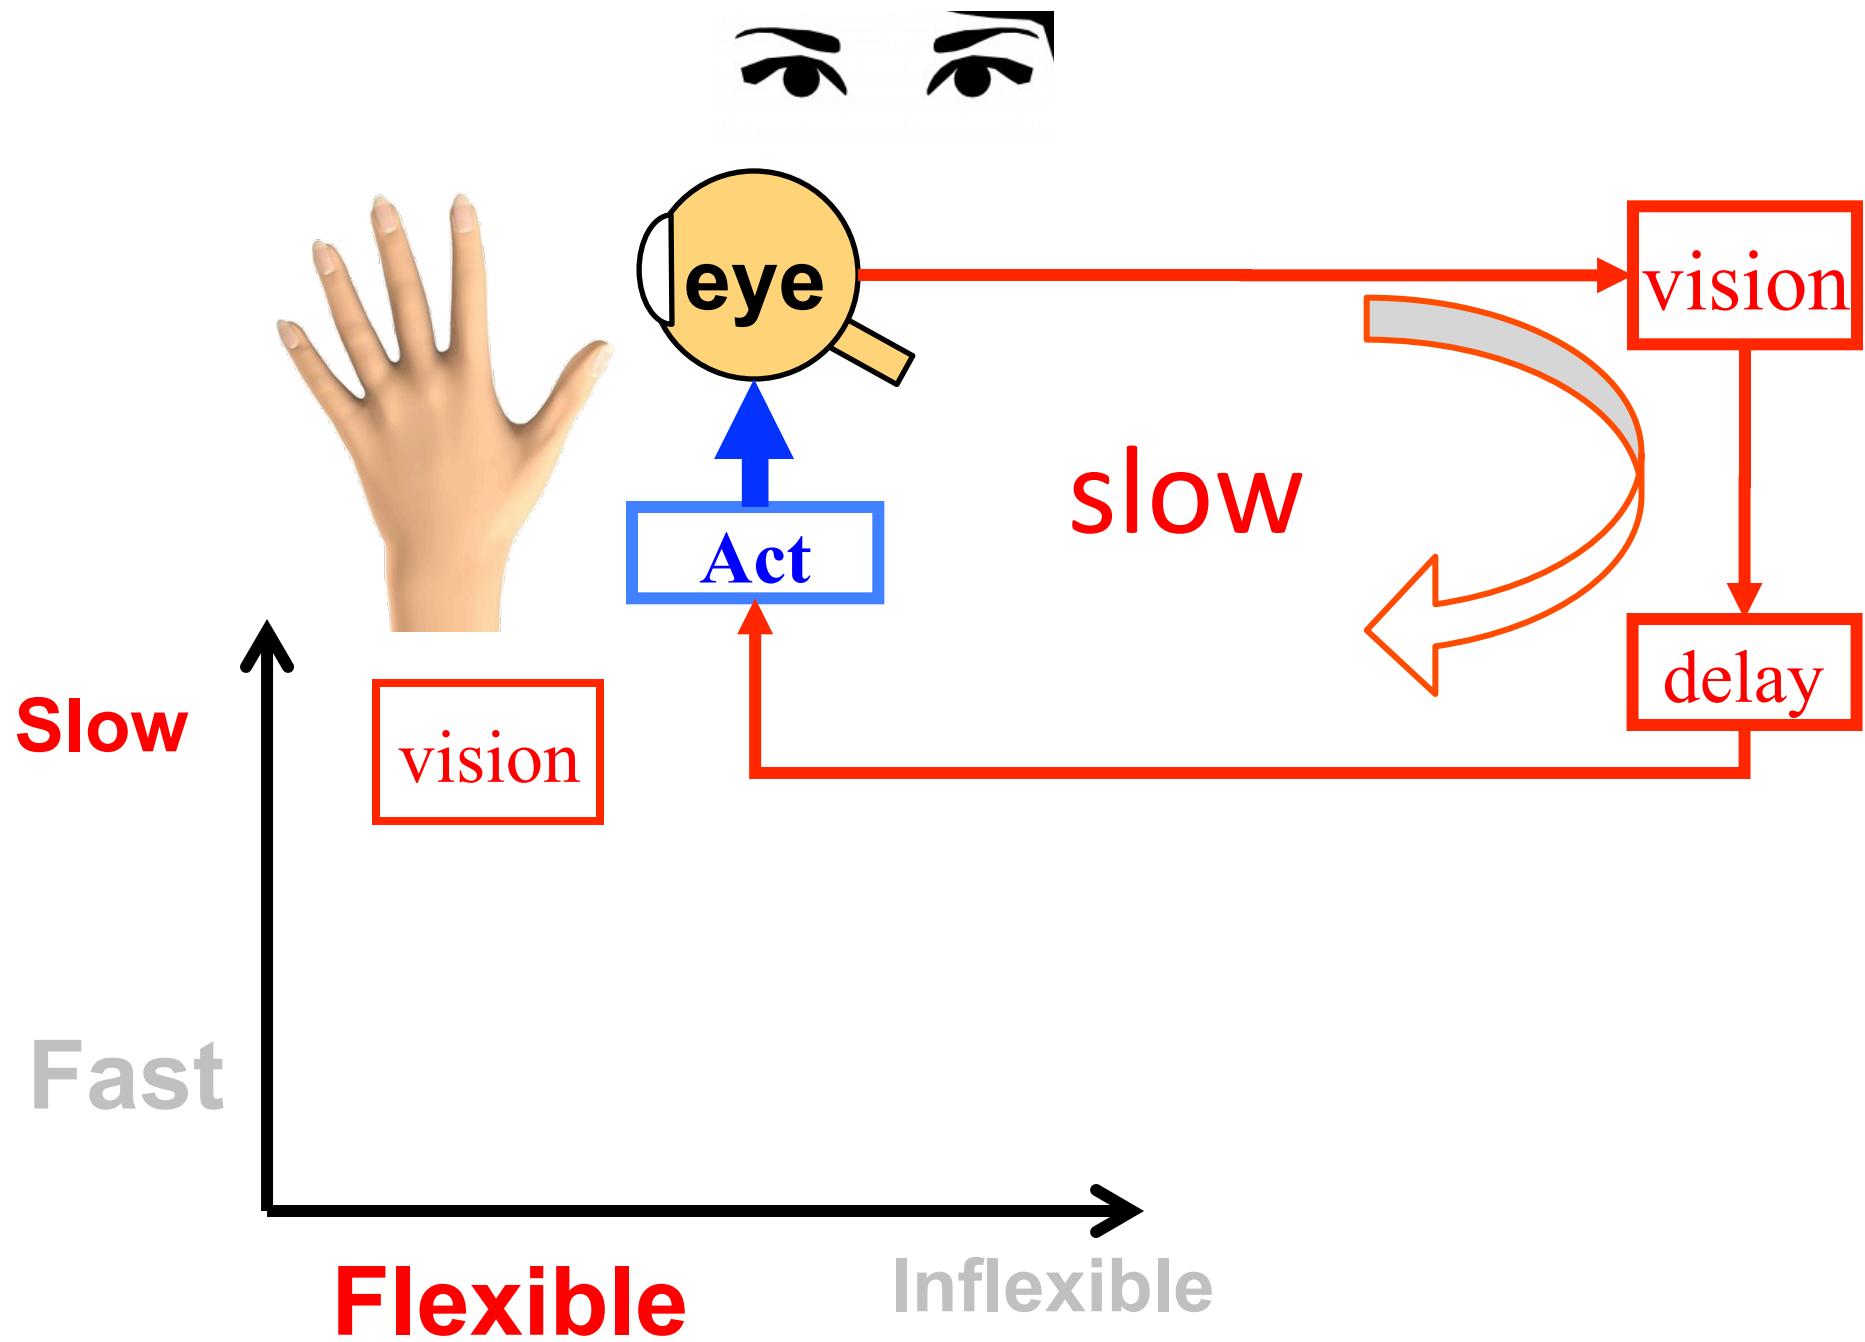
Layered Architectures and Tradeoff Spaces In Biology

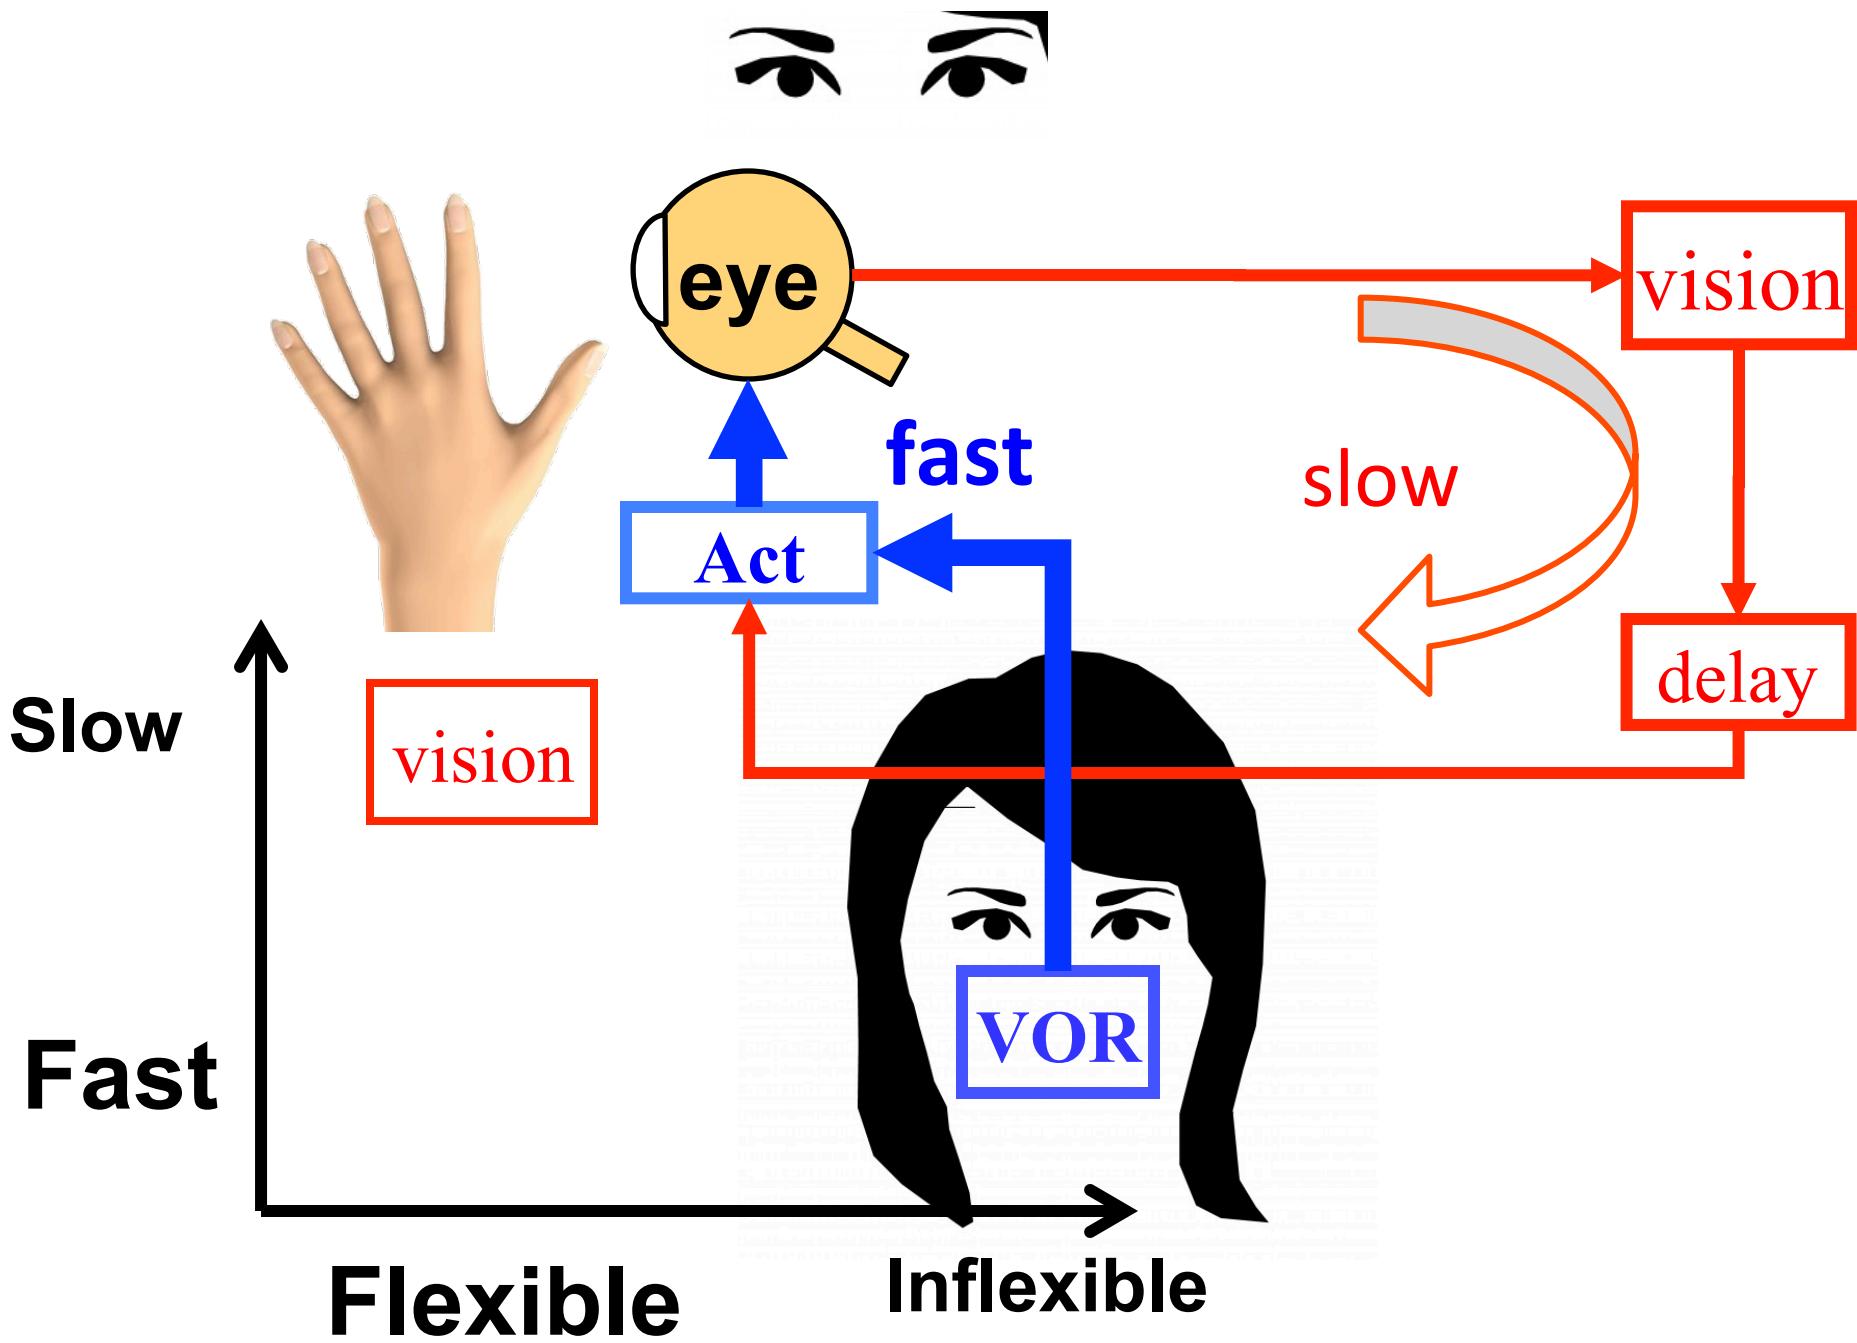


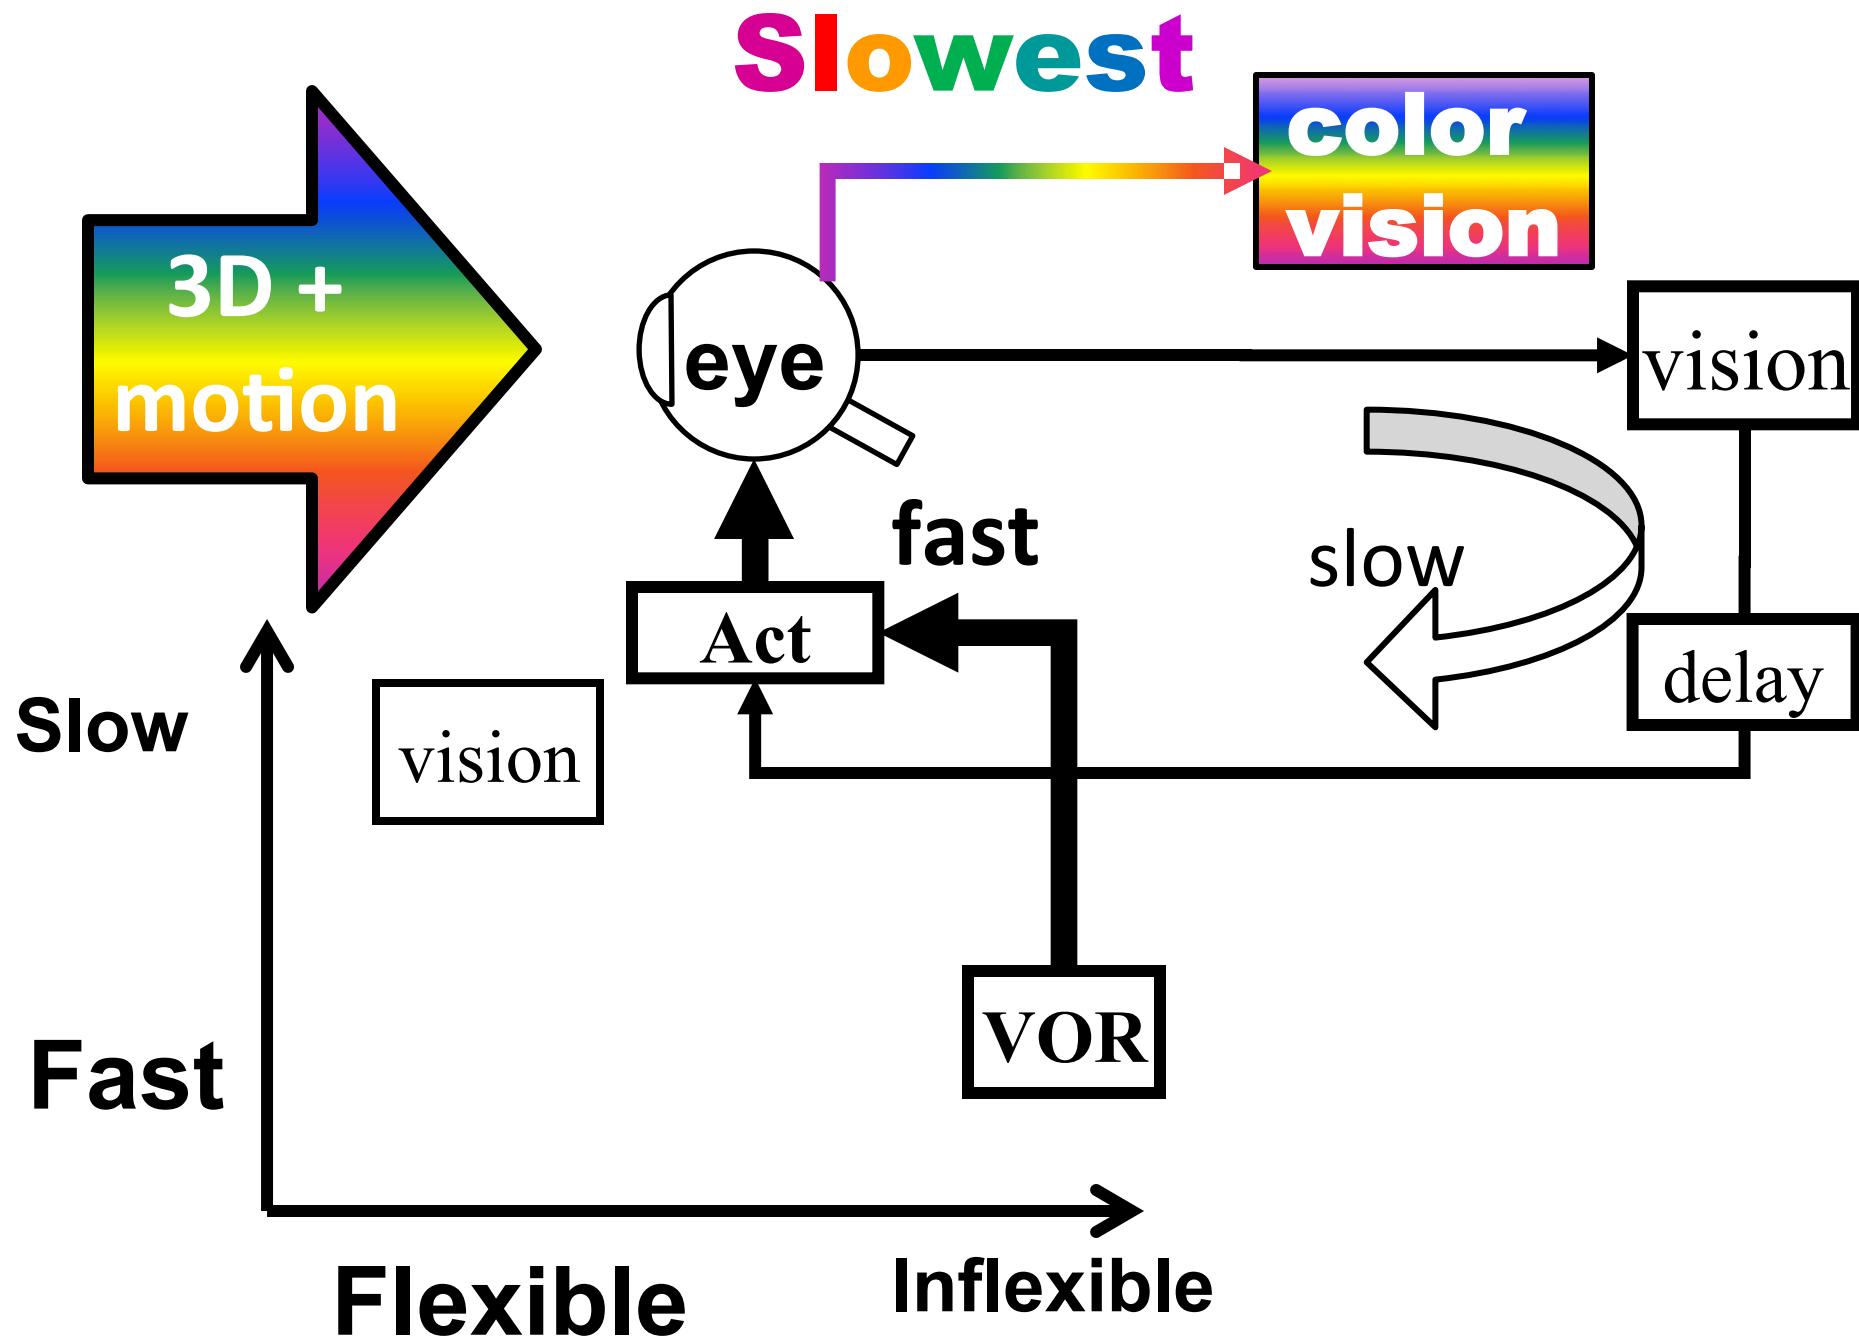
Neuroscience Example: The Vestibular Ocular Reflex (VOR)

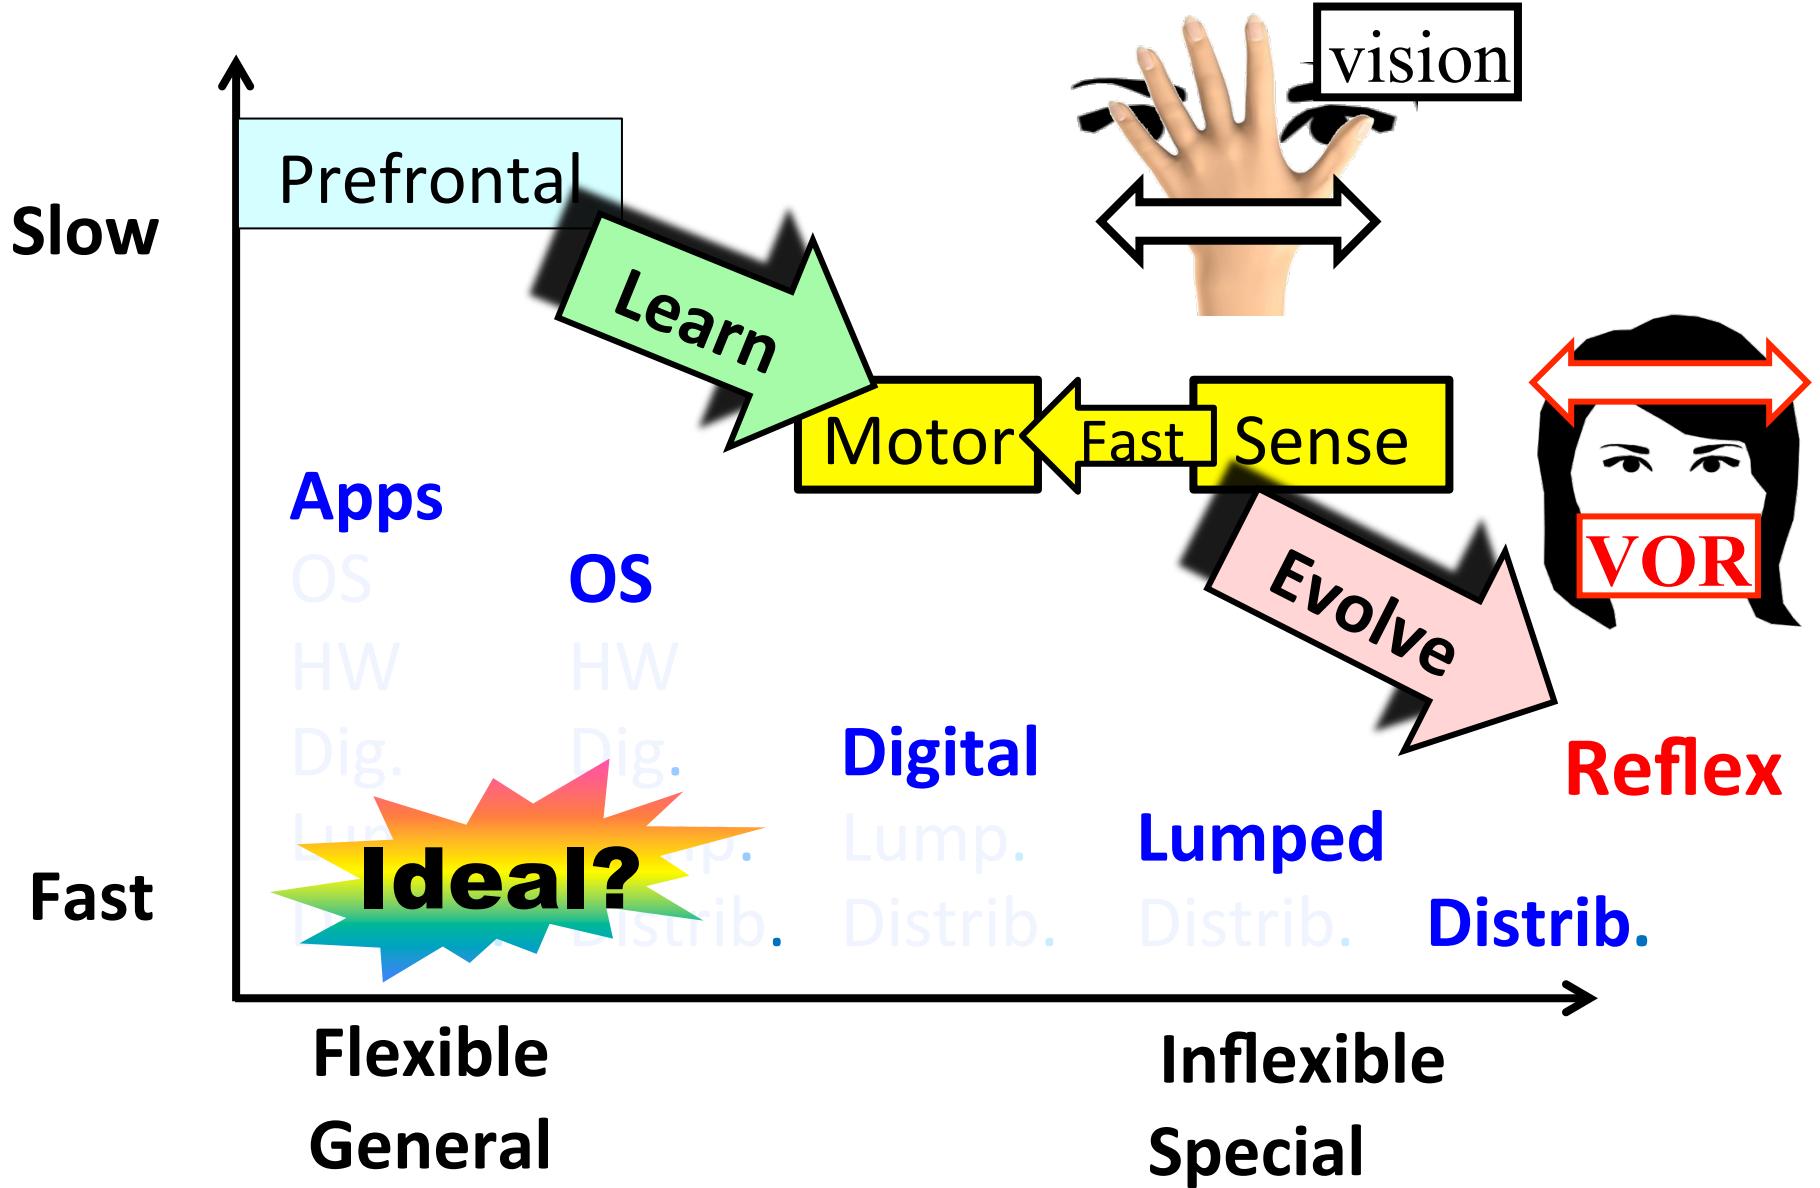
Reflex eye movement that stabilizes images on the retina during head movement











So can a new architecture beat the tradeoff?

Slow

Prefrontal

Apps

OS



Fast

Flexible

General

OS

W.
g.
imp.
strib.

Digital

Lump.
Distrib.

Lumped

Distrib.

Inflexible

Special

Learn

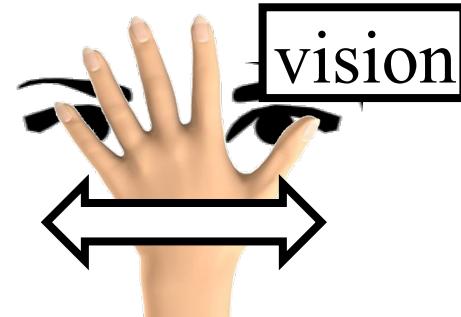
Motor

Sense

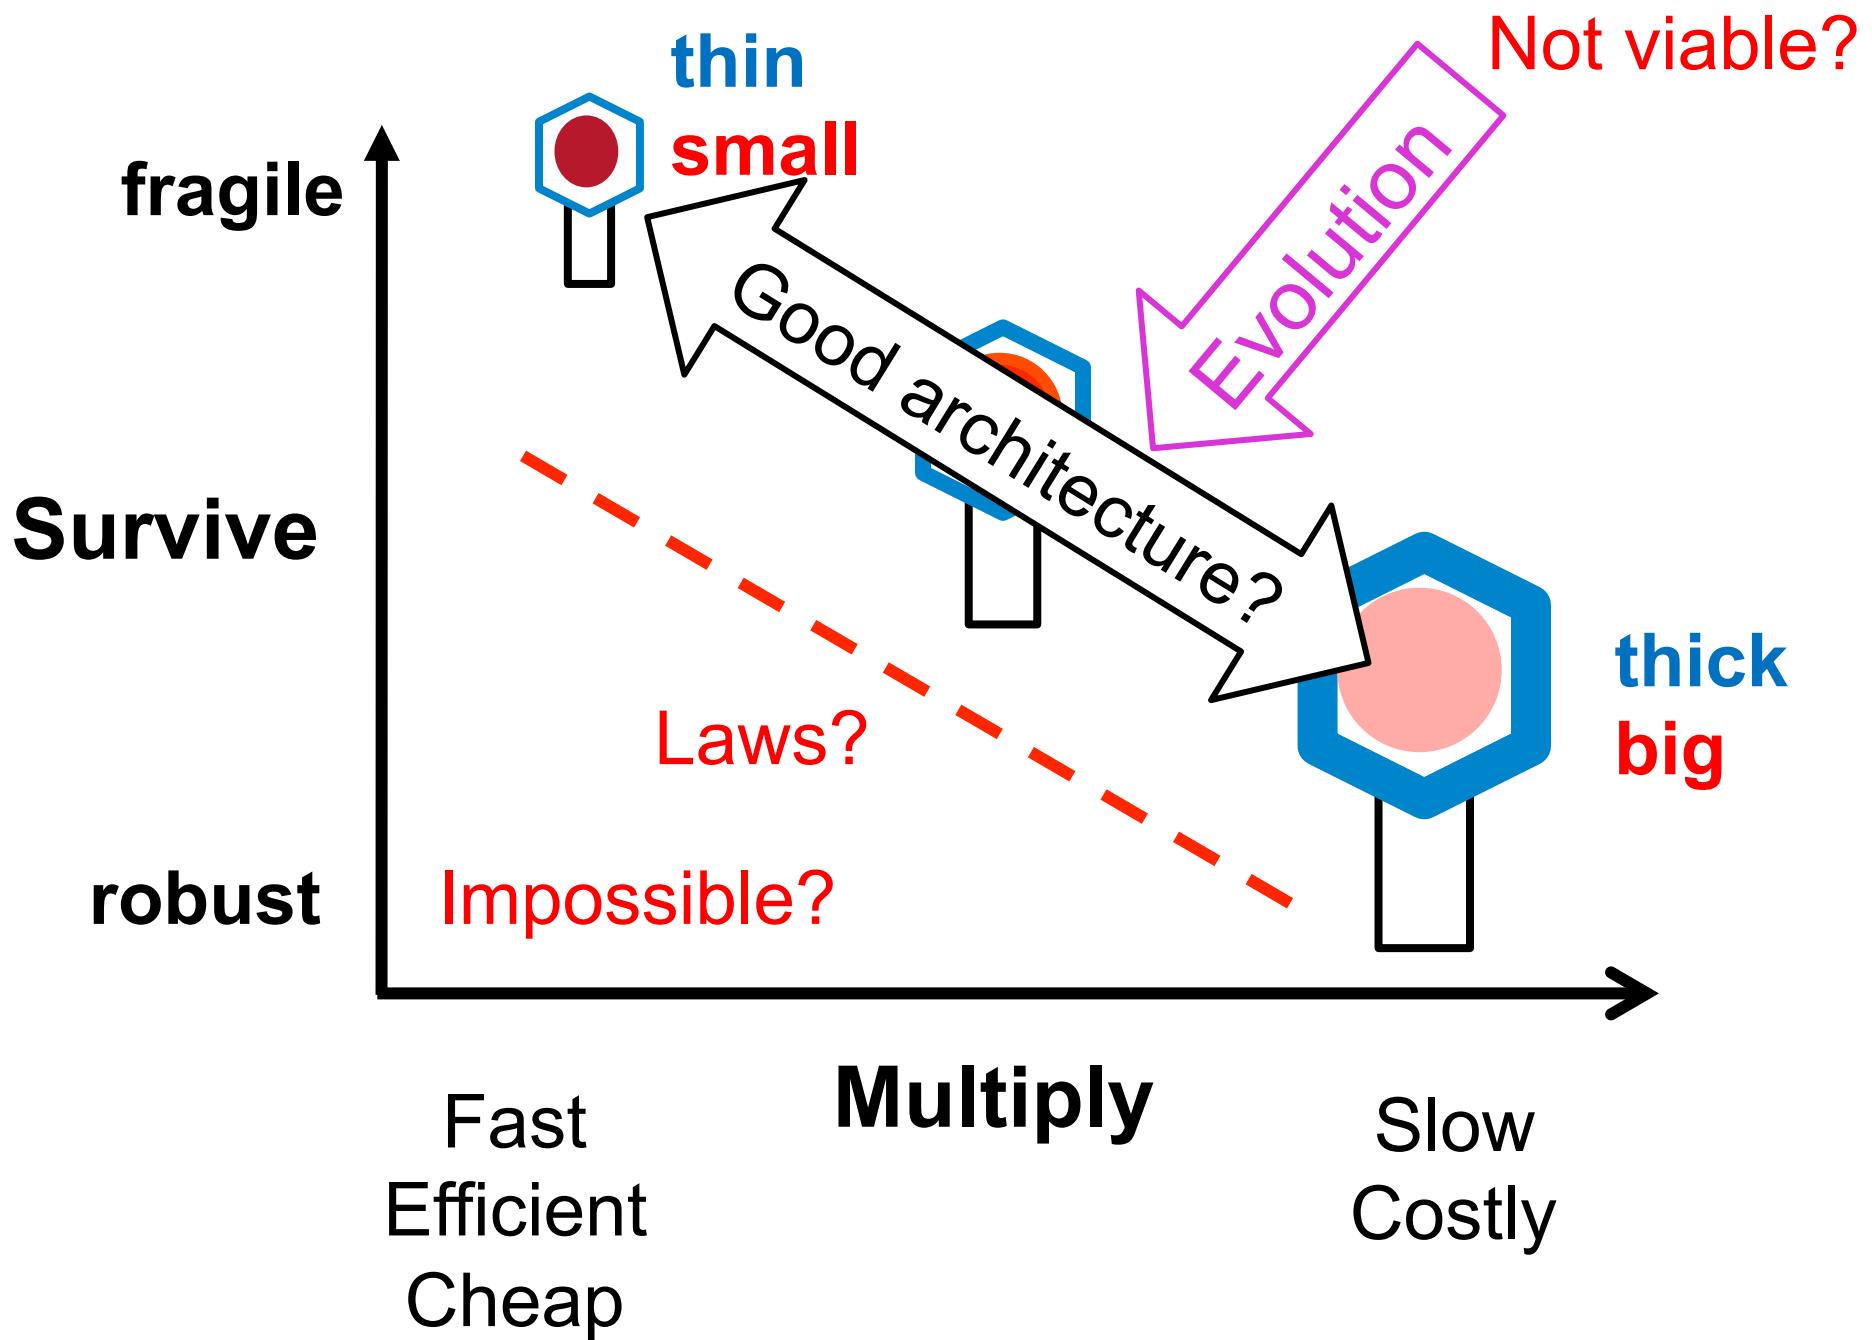
Evolve

Reflex

Distrib.



Be careful what you wish for...



Speaking of Being Careful...

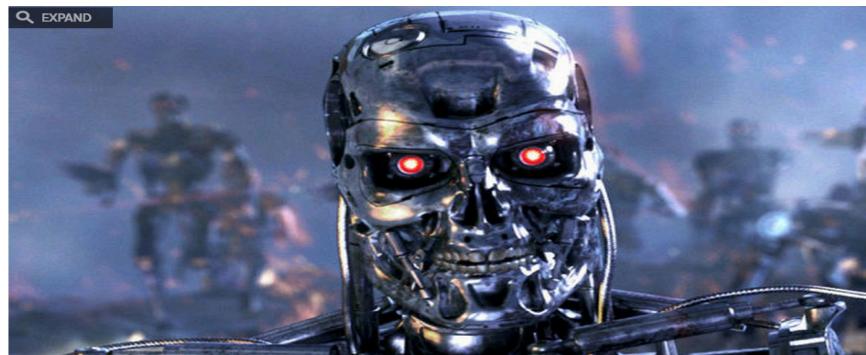
How Skynet Might Emerge From Simple Physics



George Dvorsky

Filed to: FUTURISM 4/26/13 9:00am

54,742 ⚖ 7 ★



1 2 3 4 5 ... 17

A provocative new paper is proposing that complex intelligent behavior may emerge from a fundamentally simple physical process. The theory offers novel prescriptions for how to build an AI — but it also explains how a world-dominating superintelligence might come about. We spoke to the lead author to learn more.

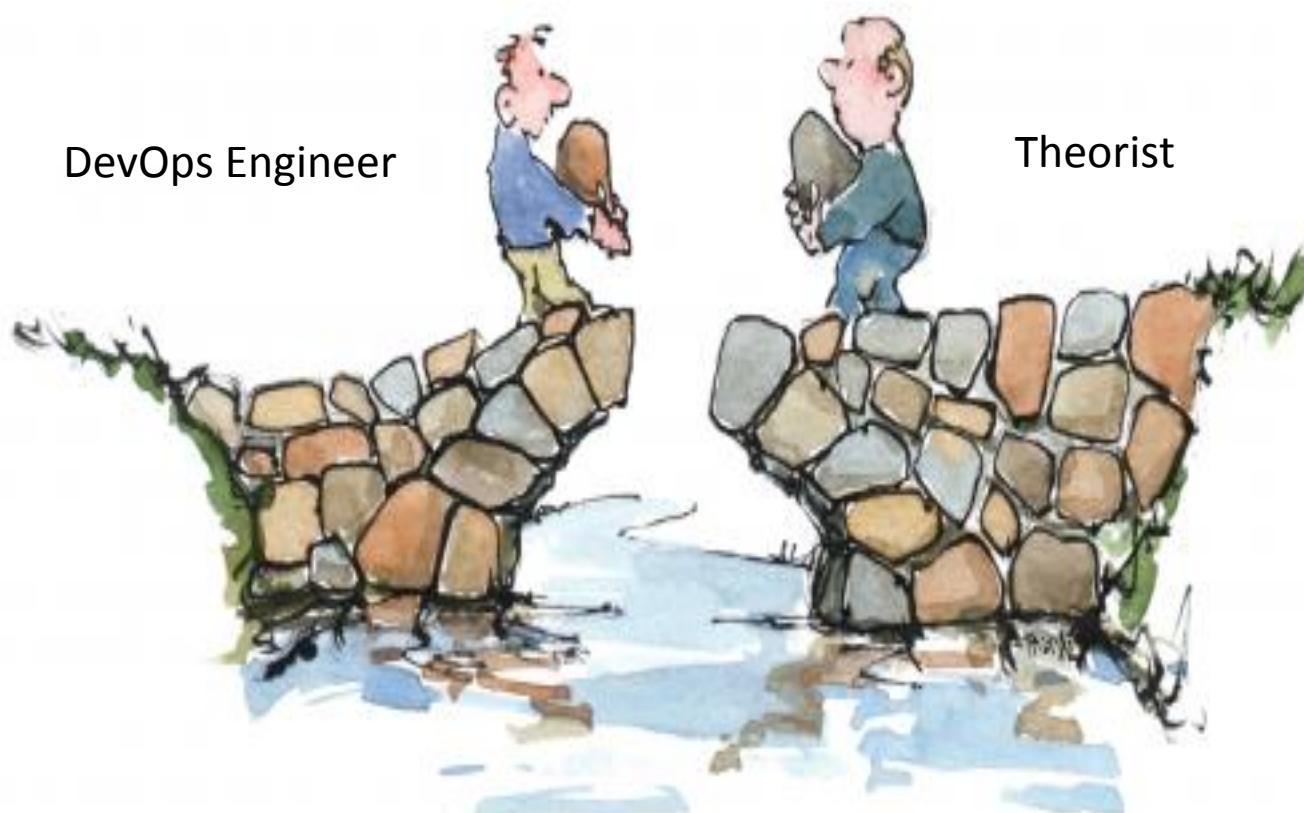
In the paper, which now appears in *Physical Review Letters*, Harvard physicist and computer scientist **Dr. Alex Wissner-Gross** posits a **Maximum Causal Entropy Production Principle** — a conjecture that intelligent behavior in general spontaneously emerges from an agent's effort to ensure its freedom of action in the future. According to this theory, intelligent systems move towards those configurations which maximize their ability to respond and adapt to future changes.

- [Dr. Alex Wissner-Gross](#) -- Maximum Causal Entropy Production Principle
- This is a conjecture that intelligent behavior in general spontaneously emerges from
 - an agent's effort to ensure its freedom of action in the future
- → Intelligent systems move towards the configurations which maximize
 - their ability to respond and adapt to future changes

Agenda

- ~~Goals for this Talk~~
- ~~What is Complexity and Why is it Hidden~~
 - ~~Robustness, Fragility, and Complexity~~
- ~~Tradeoffs?~~
- ~~The Architecture of Complex Systems~~
 - ~~Universal Architectural Principles~~
- The Evolution of DevOps
- A Few Conclusions and Q&A if we have time

Evolving DevOps



By HikingArtist.com

Evolving DevOps

- Developers + NetOps = DevOps
 - Bridging development and operations
- However, networking is a young discipline
 - Most of network design/operation/...
 - is done via engineering heuristics
 - We need deeper/better understanding of these networks to get to the next level of scale
- Building bridges?
 - *Theorists need engineers to understand what is real*
 - *Engineers need the tools theorists can provide*
 - *We need to think about how to bridge the two*
- Evolving DevOps → Bridging theory and practice
 - Theory + DevOps = <need a good name>
- One attempt
 - <http://conferences.sigcomm.org/sigcomm/2014/tutorial-theory+eng.php>

Hopefully I've Convinced You...

- That there are *Universal Architectural Features* that are common to biology and technology
- Laws, constraints, tradeoffs
 - Robust/fragile
 - Efficient/wasteful
 - Fast/slow
 - Flexible/inflexible
- Architecture/Layering
- Hidden RYF Complexity
- Hijacking, parasitism, predation
- Ok, but why is this useful?

Why is all of this Useful?

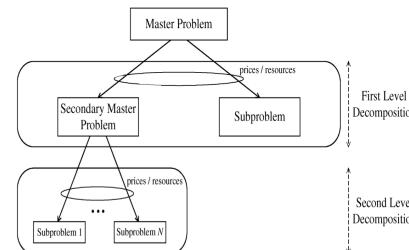
- Robust systems are intrinsically hard to understand
 - RYF is an inherent property of both advanced technology and biology
 - Understanding general principles informs what we build
 - Software (e.g., SDN, NFV, Cloud, ...) exacerbates the situation
 - And the Internet has reached an unprecedented level of complexity →
 - Need new/analytic ways of designing, deploying, and operating networks
- Nonetheless, many of our goals for the Internet architecture revolve around how to achieve robustness...
 - which requires a deep understanding of the *necessary interplay between complexity and robustness, modularity, feedback, and fragility*¹
 - which is neither accidental nor superficial
 - Rather, architecture arises from “designs” to cope with uncertainty in environment and components
 - The same “designs” make some protocols hard to evolve
 - Can anyone say, um, IPv6 (or even DNSSEC)?

¹ See Marie E. Csete and John C. Doyle, “Reverse Engineering of Biological Complexity”,
<http://www.cds.caltech.edu/~doyle/wiki/images/0/05/ScienceOnlinePDF.pdf>

Why is all of this Useful, cont?

- This much seems obvious
 - Understanding these universal architectural features and tradeoffs will help us achieve the scalability and evolvability (operability, deployability, understandability) that we are seeking from the Internet architecture today and going forward
- Perhaps less obvious: This requires a mathematical theory of network architecture
 - Want to be able to analyze/compare/simulate/optimize all aspects of network design/operation
 - Mathematics the natural language for this
 - BTW, as engineers we solve problems (“engineers always know first”), but we can benefit from the tools that theory can provide to help us design/deploy/operate/optimize our networks
- First Cut: “Layering as Optimization Decomposition: A Mathematical Theory of Network Architectures”¹
 - Network Utility Maximization (NUM)/Layering As Optimization (LAO)/Decomposition Theory

$$\begin{aligned} & \text{maximize} \quad \sum_s U_s(x_s, P_{e,s}) + \sum_j V_j(w_j) \\ & \text{subject to} \quad Rx \leq c(w, P_e), \\ & \quad x \in \mathcal{C}_1(P_e), \quad x \in \mathcal{C}_2(F) \text{ or } \in \Pi(w), \\ & \quad R \in \mathcal{R}, \quad F \in \mathcal{F}, \quad w \in \mathcal{W}. \end{aligned}$$



- TCP and Stable Path Problem (BGP) reverse-engineered as *Generalized* NUM problems
- Need something like LAO/G-NUM + “Constraints that deconstrain” view
- Evolving DevOps means bridging the DevOps and theory communities
 - <http://conferences.sigcomm.org/sigcomm/2014/tutorial-theory+eng.php>

¹ <http://www.princeton.edu/~chiangm/layering.pdf>

Q&A

Thanks!